

Computational Learning Theory

Decidability

- ▶ Computation
 - ▶ Decidability – which problems have algorithmic solutions
- ▶ Machine Learning
 - ▶ Feasibility – what assumptions must we make to trust that we can learn an unknown target function from a sample data set

Complexity

Complexity is a measure of efficiency. More efficient solutions use fewer resources.

- ▶ Computation – resources are time and space
 - ▶ Time complexity – as a function of problem size, n , how many steps must an algorithm take to solve a problem
 - ▶ Space complexity – how much memory does an algorithm need
- ▶ Machine learning – resource is data
 - ▶ Sample complexity – how many training examples, m , are needed so that with probability $\geq \delta$ we can learn a classifier with error rate lower than ϵ

Practically speaking, computational learning theory is about how much data we need to

Feasibility of Machine Learning

Machine learning is feasible if we adopt a probabilistic view of the problem and make two assumptions:

- ▶ Our training samples are drawn from the same (unknown) probability distribution as our test data, and
- ▶ Our training samples are drawn independently (with replacement)

These assumptions are known as the *i.i.d assumption* – data samples are independent and identically distributed (to the test data).

So in machine learning we use a data set of samples to make a statement about a population.

The Hoeffding Inequality

If we are trying to estimate some random variable μ by measuring ν in a sample set, the Hoeffding inequality bounds the difference between in-sample and out-of-sample error by

$$\mathbb{P}[|\nu - \mu| > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

So as the number of our training samples increases, the probability decreases that our in-sample measure ν will differ from the population parameter μ it is estimating by some error tolerance ϵ .

The Hoeffding inequality depends only on N , but this holds only for some parameter. In machine learning we are trying to estimate an entire function.

The Hoeffding Inequality in Machine Learning

In machine learning we're trying to learn an $h(\vec{x}) \in \mathcal{H}$ that approximates $f : \mathcal{X} \rightarrow \mathcal{Y}$.

- ▶ In the learning setting the measure we're trying to make a statement about is *error* and
- ▶ we want a bound on the difference between in-sample error ¹:

$$E_{in}(h) = \frac{1}{N} \sum_{n=1}^N \mathbb{I}[h(\vec{x}) \neq f(\vec{x})]$$

and out-of-sample error:

$$E_{out}(h) = \mathbb{P}[h(\vec{x}) \neq f(\vec{x})]$$

So the Hoeffding inequality becomes

$$\mathbb{P}[|E_{in}(h) - E_{out}(h)| > \epsilon] \leq 2e^{-2\epsilon^2 N}$$

But this is the error for one hypothesis.

Error of a Hypothesis Class

We need a bound for a hypothesis class. The union bound states that if $\mathcal{B}_1, \dots, \mathcal{B}_M$ are any events,

$$\mathbb{P}[\mathcal{B}_1, \text{ or } \mathcal{B}_2, \text{ or } \dots, \text{ or } \mathcal{B}_M] \leq \sum_{m=1}^M \mathbb{P}[\mathcal{B}_m]$$

For \mathcal{H} with M hypotheses h_1, \dots, h_M the union bound is:

$$\mathbb{P}[|E_{in}(g) - E_{out}(g)| > \epsilon] \leq \sum_{m=1}^M \mathbb{P}[|E_{in}(h(m)) - E_{out}(h(m))| > \epsilon]$$

If we apply the Hoeffding inequality to each of the M hypotheses we get:

$$\mathbb{P}[|E_{in}(g) - E_{out}(g)| > \epsilon] \leq 2Me^{-2\epsilon^2 N}$$

We'll return to the result later when we consider infinite hypothesis classes.

ϵ -Exhausted Version Spaces

We could use the previous result to derive a formula for N , but there is a more convenient framework based on version spaces.

Recall that a version space is the set of all hypotheses consistent with the data.

- ▶ A version space is said to be ϵ -exhausted with respect to the target function f and the data set \mathcal{D} if every hypothesis in the version space has error less than ϵ on \mathcal{D} .
- ▶ Let $|H|$ be the size of the hypothesis space.
- ▶ The probability that for a randomly chosen \mathcal{D} of size N the version space is not ϵ -exhausted is less than

$$|H|^{-\epsilon N}$$

Bounding the Error for Finite \mathcal{H}

$|H|^{-\epsilon N}$ is an upper bound on the failure rate of our hypothesis class, that is, the probability that we won't find hypothesis that has error less than ϵ on \mathcal{D} . If we want this failure rate to be no greater than some δ , then

$$|H|^{-\epsilon N} \leq \delta$$

And solving for N we get

$$N \geq \frac{1}{\epsilon} (\ln |H| + \ln \frac{1}{\delta})$$

PAC Learning for Finite \mathcal{H}

The PAC learning formula

$$N \geq \frac{1}{\epsilon} (\ln |H| + \ln \frac{1}{\delta})$$

means that we need at least N training samples to guarantee that we will learn a hypothesis that will

- ▶ *probably*, with probability $1 - \delta$ be
- ▶ *approximately*, within error ϵ
- ▶ *correct*.

Notice that N grows

- ▶ linearly in $\frac{1}{\epsilon}$,
- ▶ logarithmically in $\frac{1}{\delta}$, and
- ▶ logarithmically in $|H|$.

PAC Learning Example

Consider a hypothesis class of boolean literals. You have variables like *tall*, *glasses*, etc., and the hypothesis class represents whether a person will get a date. How many examples of people who did and did not get dates do you need to learn with 95% probability a hypothesis that has error no greater than .1

First, what's the size of the hypothesis class? For each of the variables there are three possibilities: true, false, and don't care. For example, one hypothesis for variables *tall*, *glasses*, *longHair* might be:

$$tall \wedge \neg glasses \wedge true$$

Meaning that you must be tall and not wear glasses to get a date but it doesn't matter if your hair is long.

PAC Learning Example

Since there are three values for each variable the size of the hypothesis class is

$$3^d$$

If we have 10 variables then

$$N \geq \frac{1}{\epsilon} (\ln |H| + \ln \frac{1}{\delta}) = \frac{1}{.1} (\ln 3^{10} + \ln \frac{1}{.05}) = 140$$

Dichotomies

Returning to

$$\mathbb{P}[|E_{in}(g) - E_{out}(g)| > \epsilon] \leq 2Me^{-2\epsilon^2 N}$$

Where M is the size of the hypothesis class (also sometimes written $|H|$). For infinite hypothesis classes, this won't work. What we need is an *effective* number of hypotheses.

Diversity of H is captured by idea of dichotomies. For a binary target function, there are many $h \in H$ that produce the same assignments of labels. We group these into *dichotomies*.

Effective Number of Hypotheses

Growth Function

Shattering

VC Dimension

The VC-dimension d_{VC} of a hypothesis set \mathcal{H} is the largest N for which $m_{\mathcal{H}}(N) = 2^N$.

Another way to put it: VC-dimension is the maximum number of points in a data set for which you can arrange the points in such a way that \mathcal{H} shatters those points for any labellings of the points.

VC Bound

For a confidence $\delta > 0$, the VC generalization bound is:

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \frac{4m_{\mathcal{H}}(2N)}{\delta}}$$

If we use a polynomial bound on d_{VC} :

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \left(\frac{4((2N)^{d_{VC}} - 1)}{\delta} \right)}$$

VC Bound and Sample Complexity

For an error tolerance $\epsilon > 0$ (our max acceptable difference between E_{in} and E_{out}) and a confidence $\delta > 0$, we can compute the sample complexity of an infinite hypothesis class by:

$$N \geq \frac{8}{\epsilon^2} \ln \left(\frac{4((2N)^{d_{VC}} + 1)}{\delta} \right)$$

Note that N appears on both sides, so we need to solve for N iteratively. See colt.sc for an example.

If we have a learning model with $d_{VC} = 3$ and want a generalization error at most $\epsilon = 0.1$ and a confidence of 90% ($\delta = 0.05$), we get $N = 29299$

- ▶ If we try higher values for d_{VC} , $N \approx 10000d_{VC}$, which is a gross overestimate.
- ▶ Rule of thumb: you need $10d_{VC}$ training examples to get decent generalization.

VC Bound as a Penalty for Model Complexity

You can use the VC bound to estimate the number of training samples you need, but you typically just get a data set – you're given an N .

- ▶ Question becomes: how well can we learn from the data given this data set?

If we plug values into:

$$E_{out}(g) \leq E_{in}(g) + \sqrt{\frac{8}{N} \ln \left(\frac{4((2N)^{d_{VC}} - 1)}{\delta} \right)}$$

For $N = 1000$ and $\delta = 0.1$ we get

- ▶ If $d_{VC} = 1$, error bound = 0.09
- ▶ If $d_{VC} = 2$, error bound = 0.15
- ▶ If $d_{VC} = 3$, error bound = 0.21
- ▶ If $d_{VC} = 4$, error bound = 0.27

Approximation-Generalization Tradeoff

The VC bound can be seen as a penalty for model complexity. For a more complex \mathcal{H} (larger d_{VC}), we get a larger generalization error.



- ▶ If \mathcal{H} is too simple, it may not be able to approximate f .
- ▶ If \mathcal{H} is too complex, it may not generalize well.

This tradeoff is captured in a conceptual framework called the *bias-variance* decomposition which uses squared-error to decompose the error into two terms:

$$\mathbb{E}_{\mathcal{D}} = \text{bias} + \text{var}$$

Which is a statement about a particular hypothesis class over all data sets, not just a particular data set.

Bias-Variance Tradeoff



- ▶ \mathcal{H}_1 (on the left) are lines of the form $h(x) = b$ – high bias, low variance
- ▶ \mathcal{H}_2 (on the right) are lines of the form $h(x) = ax + b$ – low bias, high variance

Total error is a sum of errors from bias and variance, and as one goes up the other goes down. Try to find the right balance. We'll learn techniques for finding this balance.