

Artificial Intelligence

Probabilistic Inference

Christopher Simpkins

Kennesaw State University

1 Approximate Inference for Bayesian Networks

Exact inference in large Bayesian networks is intractable, so we need approximate inference methods. The methods we'll learn are randomized sampling algorithms, a.k.a., **Monte Carlo** algorithms, a.k.a., **stochastic simulation**. They work by

- generating random events based on the probabilities in the Bayes net and
- counting up the different answers found in those random events.


The more the samples, the closer to the true probability distribution.

We'll learn two families of Monte carlo algorithms:

- direct sampling, and
- Markov chain sampling.

2 The Utility of Sampling

We observe data generated by some process.



Then we use those samples to make inferences about the data generating process.

- In supervised machine learning we assume a distribution for the data generating process and estimate its parameters.
- In probabilistic inference, we have a model of the full joint distribution (a Bayes net) and use it to generate samples to answer probabilistic queries.

[^AOS](#): [\[All of Statistics\]](#)

3 Computational Sampling of a Univariate Distribution

The primitive element in any sampling algorithm is the generation of samples from a known probability distribution.


To generate samples from a single discrete or real-valued variable, X ,

- First totally order the values in the domain of X .
 - For discrete variables, if there is no natural order, just create an arbitrary ordering.

Given this ordering, the cumulative probability distribution is a function of x defined by

$$f(x) = P(X \leq x).$$

- Select a uniform random number y in the domain $[0, 1]$.
- Let v be the value of X that maps to y in the cumulative probability distribution. That is, v is the element of $\text{domain}(X)$ such that
 - $f(v) = y$ or, equivalently, $v = f^{-1}(y)$
- $X = v$ is a random sample of X , chosen according to the distribution of X .



This procedure works for a single variable. If we have multiple variables and a Bayes net to represent the joint distribution over them, there are several algorithms we can use, a few of which we will see today.

4 Prior Sampling, a.k.a., Forward Sampling

To sample from a Bayes net with no associated evidence, sample each node in topological order.

- Sample unconditioned nodes according to their prior distributions.
 - Samples become fixed values for sampling CPTs of child nodes.
- Continue sampling child nodes, in order, until all variables are sampled.

function PRIOR-SAMPLE(*bn*) **returns** an event sampled from the prior specified by *bn*

inputs: *bn*, a Bayesian network specifying joint distribution $\mathbf{P}(X_1, \dots, X_n)$

$\mathbf{x} \leftarrow$ an event with n elements

for each variable X_i **in** $X_1, \dots, X_n$ **do**


$\mathbf{x}[i] \leftarrow$ a random sample from $\mathbf{P}(X_i | \text{parents}(X_i))$

return $\mathbf{x}$

5 Example: Generating a Sample from Sprinkler Bayes Net

1. Sample from $P(\textit{Cloudy}) = \langle 0.5, 0.5 \rangle$
 - Get value of true.
2. Sample from $P(\textit{Sprinkler} \mid c) = \langle 0.1, 0.9 \rangle$
 - Get value of false.
3. Sample from $P(\textit{Rain} \mid c) = \langle 0.8, 0.2 \rangle$
 - Get value of true.
4. Sample from $P(\textit{WetGrass} \mid \neg s, r) = \langle 0.9, 0.1 \rangle$
 - Get value of true.

Full sampled event: $[true, false, true, true]$



6 From Sampling to Probability Estimates

A Bayes net represents a full joint distribution, so a sample generated from a Bayes net is a sample from the prior joint distribution over all the variables. So, if S_{PS} is a sample generated by the PRIOR-SAMPLE algorithm, then:

$$S_{PS}(x_1, \dots, x_n) = \prod_{i=1}^n Pr(x_i | \text{parents}(X_i)) = P(x_1, \dots, x_n)$$

To estimate these probabilities using samples, we simply generate N samples, count the number of events of interest among the N samples, and divide that number by N . This should converge to the true probability:

$$\lim_{N \rightarrow \infty} \frac{N_{PS}(x_1, \dots, x_n)}{N} = S_{PS}(x_1, \dots, x_n) = P(x_1, \dots, x_n)$$

From the CPTs in the Bayes net, the probability of the example event we sampled, $[true, false, true, true]$ is:

$$S_{PS}(true, false, true, true) = 0.5 \cdot 0.9 \cdot 0.8 \cdot 0.9 = 0.324.$$

So if we generated $N = 1000$ samples, we would expect to see $[true, false, true, true]$ 324 times.

7 Consistent Probability Estimates

An estimate that becomes exact in the large-sample limit is called **consistent**. A consistent estimate of the probability of any partially specified event $x_1, \dots, x_m$ for $x \leq n$ is:

$$P(x_1, \dots, x_m) \approx \frac{N_{PS}(x_1, \dots, x_m)}{N} = \hat{P}(x_1, \dots, x_m) \quad (13.6)$$

That is, $\hat{P}(x_1, \dots, x_m)$ is an approximation of $P(x_1, \dots, x_m)$ that converges to the true probability as N approaches ∞

- Forward sampling generates **prior** probabilities, that is, probabilities of events prior to observing evidence.
- What we usually want is **posterior** probabilities, $P(X | e)$ – probability of X given that we have observed evidence e .

We now turn to algorithms for computing posteriors, which use prior sampling as a subroutine.

8 Rejection Sampling

Rejection sampling is a general method for producing samples from a hard-to-sample distribution given an easy-to-sample distribution. In its simplest form, it can be used to compute conditional probabilities, that is, to determine $P(X | e)$.

Let

- $\hat{P}(X | e)$ be a vector of the probabilities for each value of $x \in X$ and
- $N_{PS}(X, e)$ be a vector of sample counts for each $x \in X$ where the sample is consistent with the evidence e (the variable $\mathbf{C}$ in the algorithm).

Then:

$$\hat{P}(X | e) = \alpha N_{PS}(X, e) = \frac{N_{PS}(X, e)}{N_{PS}(e)}$$

function REJECTION-SAMPLING($X, \mathbf{e}, bn, N$) **returns** an estimate of $\mathbf{P}(X | \mathbf{e})$

inputs: X , the query variable

$\mathbf{e}$, observed values for variables $\mathbf{E}$

bn , a Bayesian network

N , the total number of samples to be generated

local variables: $\mathbf{C}$, a vector of counts for each value of X , initially zero

for $j = 1$ **to** N **do**

$\mathbf{x} \leftarrow \text{PRIOR-SAMPLE}(bn)$

if $\mathbf{x}$ is consistent with $\mathbf{e}$ **then**

$\mathbf{C}[j] \leftarrow \mathbf{C}[j] + 1$ where x_j is the value of X in $\mathbf{x}$

return NORMALIZE($\mathbf{C}$)

9 Rejection Sampling Example


Say we use prior sampling to generate 100 samples and we want to estimate $P(\text{Rain} \mid \text{Sprinkler} = \text{true})$.

- 73 samples have $\text{Sprinkler} = \text{false}$, so are **rejected**.
- Of the 27 **accepted** samples:
 - 8 samples have $\text{Rain} = \text{true}$
 - 19 samples have $\text{Rain} = \text{false}$.

Then:

$$\begin{aligned}\hat{P}(\text{Rain} \mid \text{Sprinkler} = \text{true}) &= \frac{N_{PS}(X, \mathbf{e})}{N_{PS}(\mathbf{e})} \\ &= \frac{\langle 8, 19 \rangle}{27} \\ &= \langle 0.296, 0.704 \rangle\end{aligned}$$

Calculating directly from the CPTs in the Bayes net would give us $\langle 0.3, 0.7 \rangle$.



10 Importance Sampling

The general statistical technique of **importance sampling** aims to emulate the effect of sampling from a distribution P using samples from another distribution Q , called a **proposal distribution**, whose samples are easier to obtain.

We ensure that the answers are correct in the limit by applying a correction factor $\frac{P(\mathbf{x})}{Q(\mathbf{x})}$, also known as a **weight**, to each sample $\mathbf{x}$ when counting up the samples.

Our query variable will always be one of the nonevidence variables, Z . The idea of importance sampling is to sample from $P(\mathbf{z} \mid \mathbf{e})$:

$$\hat{P}(\mathbf{z} \mid \mathbf{e}) = \frac{N_P(\mathbf{z})}{N} \approx P(\mathbf{z} \mid \mathbf{e})$$

but using an arbitrary distribution $Q(\mathbf{z})$:

$$\hat{P}(\mathbf{z} \mid \mathbf{e}) = \frac{N_Q(\mathbf{z})}{N} \frac{P(\mathbf{z} \mid \mathbf{e})}{Q(\mathbf{z})} \approx Q(\mathbf{z}) \frac{P(\mathbf{z} \mid \mathbf{e})}{Q(\mathbf{z})} = P(\mathbf{z} \mid \mathbf{e})$$

The question is: which Q to use? We want Q as close as possible to $P(\mathbf{z} \mid \mathbf{e})$ with $Q(\mathbf{z}) \neq 0$ whenever $P(\mathbf{z} \mid \mathbf{e}) \neq 0$. The most common approach is **likelihood weighting** ...

11 Likelihood Weighting

Fix the values for the evidence variables $\mathbf{E}$ and sample all the nonevidence variables in topological order, each conditioned on its parents – guarantees each sample consistent with evidence.

Let the sampling distribution be Q_{WS} and nonevidence variables be $\mathbf{Z} = \{Z_1, \dots, Z_l\}$. Then:

$$Q_{WS}(\mathbf{z}) = \prod_{i=1}^l P(z_i \mid \text{parents}(Z_i))$$

Now we need to know the weight for each sample. By general scheme for importance sampling:

$$w(\mathbf{z}) = \frac{P(\mathbf{z} \mid \mathbf{e})}{Q_{WS}(\mathbf{z})} = \alpha \frac{P(\mathbf{z}, \mathbf{e})}{Q_{WS}(\mathbf{z})}$$

where the normalizing factor $\alpha = \frac{1}{P(\mathbf{e})}$ is the same for all samples. Since $\mathbf{z}$ and $\mathbf{e}$ cover all the variables in the Bayes net, $P(\mathbf{z}, \mathbf{e})$ is the product of the conditional probabilities for the nonevidence variables times the product of the conditional probabilities for the evidence variables:

$$\begin{aligned} w(\mathbf{z}) &= \alpha \frac{P(\mathbf{z}, \mathbf{e})}{Q_{WS}(\mathbf{z})} = \alpha \frac{\prod_{i=1}^l P(z_i \mid \text{parents}(Z_i)) \prod_{i=1}^m P(e_i \mid \text{parents}(E_i))}{\prod_{i=1}^l P(z_i \mid \text{parents}(Z_i))} \\ &= \alpha \prod_{i=1}^m P(e_i \mid \text{parents}(E_i)) \end{aligned} \tag{14.9}$$

The name of this method comes from the fact that probabilities of evidence are generally called likelihoods.

12 Likelihood Weighting Algorithm

function LIKELIHOOD-WEIGHTING($X, \mathbf{e}, bn, N$) **returns** an estimate of $\mathbf{P}(X | \mathbf{e})$

inputs: X , the query variable

$\mathbf{e}$, observed values for variables $\mathbf{E}$

bn , a Bayesian network specifying joint distribution $\mathbf{P}(X_1, \dots, X_n)$

N , the total number of samples to be generated

local variables: $\mathbf{W}$, a vector of weighted counts for each value of X , initially zero

for $j = 1$ **to** N **do**

$\mathbf{x}, w \leftarrow \text{WEIGHTED-SAMPLE}(bn, \mathbf{e})$

$\mathbf{W}[j] \leftarrow \mathbf{W}[j] + w$ where x_j is the value of X in $\mathbf{x}$

return $\text{NORMALIZE}(\mathbf{W})$

function WEIGHTED-SAMPLE($bn, \mathbf{e}$) **returns** an event and a weight

$w \leftarrow 1$; $\mathbf{x} \leftarrow$ an event with n elements, with values fixed from $\mathbf{e}$

for $i = 1$ **to** n **do**

if X_i is an evidence variable with value x_{ij} in $\mathbf{e}$

then $w \leftarrow w \times P(X_i = x_{ij} | \text{parents}(X_i))$

else $\mathbf{x}[i] \leftarrow$ a random sample from $\mathbf{P}(X_i | \text{parents}(X_i))$

return $\mathbf{x}, w$

13 Likelihood Weighting Example

Given the query

- $P(\text{Rain} \mid \text{Cloudy} = \text{true}, \text{WetGrass} = \text{true})$

and the ordering

- $\text{Cloudy}, \text{Sprinkler}, \text{Rain}, \text{WetGrass},$

we initialize $w = 1.0$ and proceed as follows:

1. Cloudy is an evidence variable with value true. Therefore, we set

“ $\{$ “

$$w \leftarrow w \cdot P(c) = 0.5.$$

““

1. Sprinkler is not an evidence variable, so sample from $P(\text{Sprinkler} \mid \text{Cloudy} = \text{true}) =$


$\langle 0.1, 0.9 \rangle$; suppose this returns false.

1. Rain is not an evidence variable, so sample from $P(\text{Rain} \mid \text{Cloudy} = \text{true}) = \langle 0.8, 0.2 \rangle$; suppose this returns true.


2. WetGrass is an evidence variable with value true. Therefore, we set

“ $\{$ “
 $w \leftarrow w \cdot P(\text{WetGrass} = \text{true} \mid \text{neg } s, r) =$
 $0.5 \cdot 0.9 \cdot 0.45.$ “

Here WEIGHTED-SAMPLE returns the event $[true, false, true, true]$ with weight 0.45, which we tally under $Rain = true$.



14 Rejection vs. Importance Sampling



15 Markov Chain Monte Carlo (MCMC) Algorithms

Instead of generating each sample from scratch, MCMC algorithms generate a sample by making a random change to the preceding sample.

Think of an MCMC algorithm as being in a particular current state that specifies a value for every variable and generating a next state by making random changes to the current state.


We'll learn two MCMC algorithms:

- Gibbs sampling, and
- Metropolis-Hastings

16 Gibbs Sampling and Markov Blankets



17 Conditional Independence Relations



- (a) Each variable is conditionally independent of its non-descendants, given its parents.
- (b) A variable is conditionally independent of all other nodes in the network, given its parents, children, and children's parents – that is, given its **Markov blanket**.

For example, the variable *Burglary* is independent of *JohnCalls* and *MaryCalls*, given *Alarm* and *Earthquake*.

18 Gibbs Sampling


- Fix the evidence variables.
- Start in a arbitrary state sampled from the nonevidence variables.
 - Each X_i is sampled from the values in its Markov blanket, $mb(X_i)$

“{=latex}


$$P(x_i \mid mb(X_i)) = \alpha P(x_i \mid \text{parents}(X_i)) \prod_{Y_j \in \text{Children}(X_i)} P(y_j \mid \text{parents}(Y_j)) \quad (13.10)$$

““

- Wander the state space randomly, keeping the evidence variables fixed and flipping one nonevidence variable by sampling from a distribution calculated Equation 13.10.



(a)



(b)


Gibbs sampling for X_i means sampling conditioned on the current values of the variables in its Markov blanket.

19 Gibbs Sampling Example: Step 1

Consider the query $P(\text{Rain} \mid \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$.

- Evidence variables *Sprinkler* and *WetGrass* fixed to observed values (true), nonevidence variables *Cloudy* and *Rain* initialized randomly to, say, true and false respectively.
 - Initial state is $[\text{true}, \mathbf{\text{true}}, \text{false}, \mathbf{\text{true}}]$. Fixed evidence variables marked in bold.

Now the nonevidence variables Z_i are sampled repeatedly in some random order according to a probability distribution $\rho(i)$ for choosing variables. For example:



1. *Cloudy* is chosen and then sampled, given the current values of its Markov blanket: in this case, we sample from $P(\text{Cloudy} | \text{Sprinkler} = \text{true}, \text{Rain} = \text{false})$ whose distribution is calculated using Equation 13.10:

$$P(x_i | mb(X_i)) = \alpha P(x_i | \text{parents}(X_i)) \prod_{Y_j \in \text{Children}(X_i)} P(y_j | \text{parents}(Y_j)) \quad (13.10)$$

$$P(c | s, \neg r) = \alpha P(c) P(s | c) P(\neg r | c) = \alpha 0.5 \cdot 0.1 \cdot 0.2$$

$$P(\neg c | s, \neg r) = \alpha P(\neg c) P(s | \neg c) P(\neg r | \neg c) = \alpha 0.5 \cdot 0.5 \cdot 0.8$$

yielding $\alpha \langle 0.001, 0.020 \rangle \approx \langle 0.048, 0.952 \rangle$


- Suppose the result is *Cloudy* = *false*. Then the new current state is [*false*, **true**, *false*, **true**].

20 Gibbs Sampling Example: Step 2

1. Sampling $\rho(i)$ then chooses *Rain* as next variable. We sample from $P(\text{Rain} | \text{Cloudy} = \text{false}, \text{Sprinkler} = \text{true}, \text{WetGrass} = \text{true})$ using Equation 13.10 (just like Step 1, not shown here).
2. Suppose this yields *Rain* = *true*. The new current state is [*false*, **true**, *true*, **true**].

Figure on the right shows the complete Markov chain for the case where variables are chosen uniformly, i.e., $\rho(\text{Cloudy}) = \rho(\text{Rain}) = 0.5$.

- The algorithm wanders around in this graph, following links with stated probabilities.
- Each state visited during this process is a sample that contributes to the estimate for the query variable *Rain*. - If the process visits 20 states where *Rain* is true and 60 states where *Rain* is false, then the answer to the query is $\text{NORMALIZE}(\langle 20, 60 \rangle) = \langle 0.25, 0.75 \rangle$.



21 Gibbs Sampling Algorithm

function GIBBS-ASK($X, \mathbf{e}, bn, N$) **returns** an estimate of $\mathbf{P}(X | \mathbf{e})$

local variables: $\mathbf{C}$, a vector of counts for each value of X , initially zero
 $\mathbf{Z}$, the nonevidence variables in bn
 $\mathbf{x}$, the current state of the network, initialized from $\mathbf{e}$

initialize $\mathbf{x}$ with random values for the variables in $\mathbf{Z}$

for $k = 1$ **to** N **do**

choose any variable Z_i from $\mathbf{Z}$ according to any distribution $\rho(i)$


 set the value of Z_i in $\mathbf{x}$ by sampling from $\mathbf{P}(Z_i | mb(Z_i))$

$\mathbf{C}[j] \leftarrow \mathbf{C}[j] + 1$ where x_j is the value of X in $\mathbf{x}$

return NORMALIZE($\mathbf{C}$)

22 Gibbs Sampling vs. Importance Sampling


Performance of Gibbs sampling compared to likelihood weighting on the car insurance network.



- (a) for the standard query on *PropertyCost*, and
- (b) for the case where the output variables are observed and *Age* is the query variable.

Gibbs sampling typically outperforms likelihood weighting when evidence is mostly downstream.

23 Markov Chains



- (a) The states and transition probabilities of the Markov chain for the query $P(Rain|Sprinkler = true, WetGrass = true)$. Note the self-loops: the state stays the same when either variable is chosen and then resamples the same value it already has.
- (b) The transition probabilities when the CPT for *Rain* constrains it to have the same value as *Cloudy*.

24 Metropolis



25 Metropolis-Hastings Sampling

Like simulated annealing, MH has two stages in each iteration of the sampling process:

1. Sample a new state x' from a **proposal distribution** $q(x'|x)$, given the current state x .
2. Probabilistically accept or reject x' according to the **acceptance probability**

$$a(x'|x) = \min \left(1, \frac{\pi(x')q(x|x')}{\pi(x)q(x'|x)} \right)$$

If the proposal is rejected, the state remains at x .

$q(x'|x)$ could be defined as follows:

- With probability 0.95, perform a Gibbs sampling step to generate x' .
- Otherwise, generate x' by running WEIGHTED-SAMPLE using likelihood weighting.

This has the effect of getting MH "unstuck."