

Relational Model

CS 6070 Databases

Kennesaw State University

Relational Data Model

A *relation schema* $R(A_1, \dots, A_n)$ is a relation name R and a list of attributes A_1, \dots, A_n .

Each attribute A_i is the name of a role played by some domain D .

- ▶ Example: $AUTHOR(author_id, first_name, last_name)$
 - ▶ $dom(A_1)$ (or $dom(author_id)$) is integer
 - ▶ The role played by an integer in A_1 is that of an identifier/key.

A *database schema* is a collection of relation schemas.

- ▶ Example: *PUBS* database has relation schemas $BOOK$, $AUTHOR$, and PUB (for publication, not public house)

Relations and Databases

A *relation*, or *relation state*, $r(R)$ is a **set** of tuples that conform to a *relation schema* R .

- ▶ Example: given $AUTHOR(author_id, first_name, last_name)$, a particular $r(AUTHOR) =$

author_id	first_name	last_name
1	John	McCarthy
4	Claude	Shannon
5	Alan	Turing
6	Alonzo	Church

A *database* is a set of relations.

Tuples

A *tuple* is an **ordered list** of values.

- ▶ Example: $t_1 = \langle 1, \text{'John'}, \text{'McCarthy'} \rangle$

Each value in the tuple is that tuple's value for the corresponding attribute of the relation schema.

Example: (these are equivalent notations):

- ▶ $t_1[\text{first_name}] = \text{'John'}$ (bracket notation)
- ▶ $t_1.\text{first_name} = \text{'John'}$ (object notation)
- ▶ $t_1[2] = \text{'John'}$ (positional notation)

The *degree* or *arity* of a relation schema is the number of attributes it has.

- ▶ Example: *AUTHOR* has degree 3.

Attributes and Domains

Each attribute has a name and a *domain*

- ▶ The name describes the role played by the attribute
 - ▶ Example: the *first_name* attribute of the *AUTHOR* schema plays the role of the first name of an author represented by a tuple in a $r(AUTHOR)$ relation.
- ▶ The domain is a set of atomic values that a tuple may have for that attribute.
- ▶ A domain has a *logical definition*, e.g., integer or string, and may also have a *format*.
 - ▶ Example: *Home_phone* as $ddd - dddd$, where d is a digit

The diagram illustrates the components of a relation. At the top, 'Relation Name' points to 'STUDENT'. Below it, 'Attributes' points to a table header with columns: Name, Ssn, Home_phone, Address, Office_phone, Age, and Gpa. The table body contains five tuples with data for each column. On the left, 'Tuples' points to the rows of the table.

Name	Ssn	Home_phone	Address	Office_phone	Age	Gpa
Benjamin Bayer	305-61-2435	373-1616	2918 Bluebonnet Lane	NULL	19	3.21
Chung-cha Kim	381-62-1245	375-4409	125 Kirby Road	NULL	18	2.89
Dick Davidson	422-11-2320	NULL	3452 Elgin Road	749-1253	25	3.53
Rohan Panchal	489-22-1100	376-9821	265 Lark Lane	749-6492	28	3.93
Barbara Benson	533-69-1238	839-8461	7384 Fontana Lane	NULL	19	3.25

Figure 5.1
The attributes and tuples of a relation STUDENT.

Mathematical Definition of Relation

Given $R(A_1, \dots, A_n)$,

- ▶ $r(R) \subseteq (\text{dom}(A_1) \times \text{dom}(A_2) \times \dots \times \text{dom}(A_n))$

The total number of values, or *cardinality*, of a domain D is $|D|$.

So the maximum number of tuples that could possibly be in $r(R)$ is

- ▶ $|\text{dom}(A_1)| * |\text{dom}(A_2)| * \dots * |\text{dom}(A_n)|$

Example: given

- ▶ $R(A_1, A_2)$
- ▶ $\text{dom}(A_1) = \{1, 2\}$, $\text{dom}(A_2) = \{a, b\}$

What are all the possible tuples that could appear in any $r(R)$?

Enumerating Tuples

Example: given

- ▶ $R(A_1, A_2)$
- ▶ $\text{dom}(A_1) = \{1, 2\}$, $\text{dom}(A_2) = \{a, b\}$

What are all the possible tuples that could appear in any $r(R)$?

$$\text{dom}(A_1) \times \text{dom}(A_2) = \{\langle 1, a \rangle, \langle 1, b \rangle, \langle 2, a \rangle, \langle 2, b \rangle\}$$

Given the definition of a *relation* or *relation state*, what is the maximum size of any $r(R)$?

Properties of Relations

- ▶ Atomicity of values, i.e., the First Normal Form assumption
 - ▶ Attribute values in tuples are indivisible, e.g., no compound or multivalued attributes as in EER models
- ▶ Nulls may appear in tuples for *some* attributes (more later)
 - ▶ Unknown, not applicable, not existing
- ▶ Closed world assumption
 - ▶ Facts not asserted explicitly are assumed to be false

Consider the properties above in the context of the following relation.

author_id	first_name	middle_name	last_name
1	John	NULL	McCarthy
4	Claude	Elwood	Shannon
5	Alan	Mathison	Turing
6	Alonzo	NULL	Church

Superkeys

A *superkey* SK is a set of attributes of a relation schema R such that

$$t_i[SK] \neq t_j[SK]$$

for any $i \neq j$.

In other words, the values of the superkey attributes of a tuple uniquely identify the tuple within the relation.

By the definition of the relational model, the full attribute set of a relation schema is a *default superkey*.

- ▶ Pause for a moment and make sure you understand that last statement.

Keys

A *minimal superkey* is a superkey for which removing an attribute would make it no longer a superkey.

We call a minimal superkey a *key*.

A relation schema may have several keys. We call these *candidate keys* and choose one arbitrarily to be the *primary key*.

We underline the primary key in a relation schema.

- ▶ Example: *AUTHOR*(author_id, first_name, last_name)

Database Integrity Constraints

- ▶ Inherent model-based (or *implicit*) constraints
 - ▶ Domain constraints: attribute values in tuples must be in domain for that attribute and conform to formats
 - ▶ Atomic attribute values
- ▶ Schema-based (or *explicit*) constraints
 - ▶ Key constraints: no two tuples can have the same values for the primary key
 - ▶ Entity Integrity Constraints: no tuple can have a NULL value for its primary key attribute
 - ▶ Referential Integrity Constraints: tuples in one relation referencing tuples in another relation
- ▶ Application-based (or semantic constraints), a.k.a., business rules
 - ▶ Constraints on attribute values that cannot be expressed in the relational model

Referential Integrity Constraints

A foreign key value from a tuple in one relation must refer to nothing, or to the primary key for an existing tuple in another relation. Formally:

Given relation schemas R_1 and R_2 , a set of attributes FK in R_1 is a foreign key referencing R_2 if

- ▶ the attributes in FK in R_1 have same domains as PK in R_2 , and
- ▶ given some t_1 in $r_1(R_1)$ and t_2 in $r_2(R_2)$, either $t_1[FK] = t_2[PK]$ or $t_1[FK]$ is NULL.

R_1 is the referencing, or child relation, R_2 is the referenced, or parent relation.

Diagramming FK Relationships

EMPLOYEE

Fname	Minit	Lname	<u>Ssn</u>	Bdate	Address	Sex	Salary	Super_ssn	Dno
-------	-------	-------	------------	-------	---------	-----	--------	-----------	-----



DEPARTMENT

Dname	<u>Dnumber</u>	Mgr_ssn	Mgr_start_date
-------	----------------	---------	----------------



DEPT_LOCATIONS

<u>Dnumber</u>	Dlocation
----------------	-----------

PROJECT

Pname	<u>Pnumber</u>	Plocation	Dnum
-------	----------------	-----------	------



WORKS_ON

Essn	Pno	Hours
------	-----	-------

DEPENDENT

Essn	Dependent_name	Sex	Bdate	Relationship
------	----------------	-----	-------	--------------

Semantic Integrity Constraints

- ▶ Can't be specified in DDL
- ▶ Can be checked with triggers and assertions
- ▶ Usually checked in application code

Example: salary of an employee cannot exceed the salary of the employee's supervisor.

Constraint Violations on Insert, Update or Delete

- ▶ Domain constraints
 - ▶ Insert/update a tuple with an attribute value not in attribute's domain
- ▶ Key constraints
 - ▶ Insert/update a tuple with a key that's already in the relation state
- ▶ Entity integrity constraints
 - ▶ Insert/update a tuple with a NULL value for any part of the primary key
- ▶ Referential integrity constraints
 - ▶ Insert/update a tuple in a referring relation whose FK does not appear as a PK value in any tuple of the referenced relation
 - ▶ Update the primary key for a tuple in a referenced relation for which there are tuples in referring relationships. The tuples in referring relationships would be orphaned or end up referring to the wrong parent tuple.
 - ▶ Delete a tuple in a referenced relation for which there are tuples in referring relations. The tuples in referring relations would be orphaned.

Domain Integrity Violation Examples

<i>author_id</i>	<i>first_name</i>	<i>last_name</i>
1	John	McCarthy
4	Claude	Shannon
5	Alan	Turing
6	Alonzo	Church

$\text{dom}(\text{author_id}) = \text{integer}$, $\text{dom}(\text{first_name}) = \text{string}$, $\text{dom}(\text{last_name}) = \text{string}$

- ▶ Insert <"Two", "Jenny", "McCarthy"> – "Two" is not in $\text{dom}(\text{author_id})$
- ▶ Update <1, "John", "McCarthy"> to <1, "John", 1> – 1 is not in $\text{dom}(\text{last_name})$

Key Integrity Violation Examples

<i>author_id</i>	<i>first_name</i>	<i>last_name</i>
1	John	McCarthy
4	Claude	Shannon
5	Alan	Turing
6	Alonzo	Church

- ▶ Insert <1, "Jenny", "McCarthy"> – 1 is an existing primary key
- ▶ Update <1, "John", "McCarthy"> to <6, "John", "McCarthy"> – 6 is an existing primary key

Entity Integrity Violation Examples

<i>author_id</i>	<i>first_name</i>	<i>last_name</i>
1	John	McCarthy
4	Claude	Shannon
5	Alan	Turing
6	Alonzo	Church

- ▶ Insert <NULL, "Jenny", "McCarthy"> — NULL not allowed for primary key
- ▶ Update <NULL, "John", "McCarthy"> to <1, "John", 1>— NULL not allowed for primary key

Referential Integrity Violation Examples

author

<u>id</u>	first_name	last_name
1	John	McCarthy
4	Claude	Shannon
5	Alan	Turing
6	Alonzo	Church

`pub[author_id]` is a foreign key to the `author` relation.

pub

<u>id</u>	title	author_id
1	Recursive Functions of Symbolic ...	1
2	A Mathematical Theory of Commun- ication	4
4	Computing machinery and intelli- gence	5
5	The calculi of lambda-conversion	6

- ▶ Update `<1, "John", "McCarthy">` to `<2, "John", "McCarthy">` – breaks reference from `pub[id] = 1`.
- ▶ Delete `<5 "Alan", "Turning">` from `author` – orphans tuple from `pub`.
- ▶ Update `<5, "The calculi of lambda-conversion", 6>` – `<5, "The calculi of lambda-conversion", 7>` – 7 is not a primary key in `pub`.
- ▶ Update `<5, "The calculi of lambda-conversion", 6>` – `<5, "The calculi of lambda-conversion", NULL>` – strictly speaking, not a violation, but creates bad data. We'll see how to prevent that when we learn SQL.

Closing Thoughts

- ▶ The relational model is a mathematical database model.
- ▶ Aside from its rigorous grounding, a strength of the relational model is its modeling (and, in DBMS systems, enforcement) of constraints.
- ▶ Relational data model is especially valuable if the integrity of your data is important.