

Probability, etc.

CS 4277 Deep Learning

Kennesaw State University

Probability¹

Probability theory: quantification and manipulation of uncertainty.

- ▶ Epistemic, a.k.a. systematic uncertainty: we only see data sets of finite size
- ▶ Aleatoric, a.k.a. intrinsic, stochastic uncertainty: noise – we only observe partial information




60%



40%

Probability in Machine Learning

We observe data generated by a random process.




2

We make some assumptions about the data generating function and infer its parameters using samples from the process (training data).

A Medical Screening Example

A cancer with occurrence rate of 1% (.01) has a “90% accurate” test, and:



False positive rate: .03, False negative rate: 0.10

Questions:

- ▶ If we screen someone, what is the probability that they test positive?
- ▶ If someone tests positive, what is the probability that they have cancer?

We'll return to these questions after we develop some analysis tools.

Joint Probability


Let X and Y be *random* (a.k.a. *stochastic*) variables and

- ▶ $\{x_i\}_{i=1}^L$
- ▶ $\{y_j\}_{j=1}^M$
- ▶ N trials in which we sample X and Y
- ▶ n_{ij} is number of trials in which $X = x_i$ and $Y = y_j$
- ▶ c_i is the number of trials in which $X = x_i$, for all y s
- ▶ r_j is the number of trials in which $Y = y_j$, for all x s

Then the joint probability of observing x_i and y_j is

$$p(X = x_i, Y = y_j) = \frac{n_{ij}}{N}$$

We can visualize this event with the grid diagram on the right. Note that we're always observing events where both random variables have values, e.g., when we screen a person for cancer we're observing a joint event of two random variables: the test result and the actual existence of cancer.



The Sum Rule

$$p(X = x_i) = \frac{c_i}{N}$$

Notice that the number of instances in column i , c_i , is the sum of instances in each cell having n_{ij} instances, so $c_i = \sum_j n_{ij}$. Recalling that

$$p(X = x_i, Y = y_j) = \frac{n_{ij}}{N}$$

we have

$$p(X = x_i) = \sum_{j=1}^M p(X = x_i, Y = y_j)$$

This is the *sum rule*, which is also called the marginal probability because we sum over the other variable and write the sum in the margin of the table.

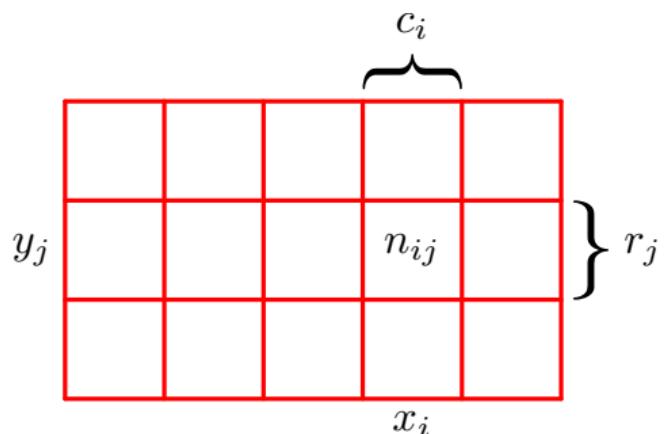
Conditional Probability

If we consider trials in which $X = x_i$, the fraction of those trials in which $Y = y_j$ is written

$$p(Y = y_j | X = x_i)$$

We call this the *conditional probability* of $Y = y_j$ given $X = x_i$, which is the fraction of points in column i that fall in cell i, j so:

$$p(Y = y_j | X = x_i) = \frac{n_{ij}}{c_i}$$



The Product Rule

Given the previous definitions for conditional probabilities and marginal probabilities, we can derive a formula for joint probabilities:

$$p(X = x_i, Y = y_j) = \frac{n_{ij}}{N} = \frac{n_{ij}}{c_i} \frac{c_i}{N}$$
$$p(X = x_i, Y = y_j) = p(Y = y_j | X = x_i)p(X = x_i)$$

This is the *product rule*.


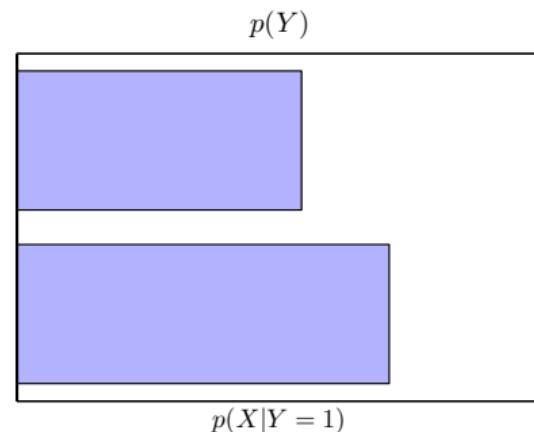
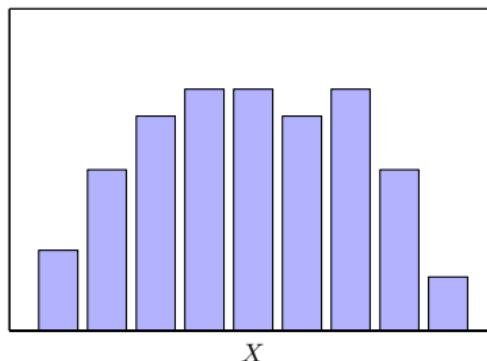
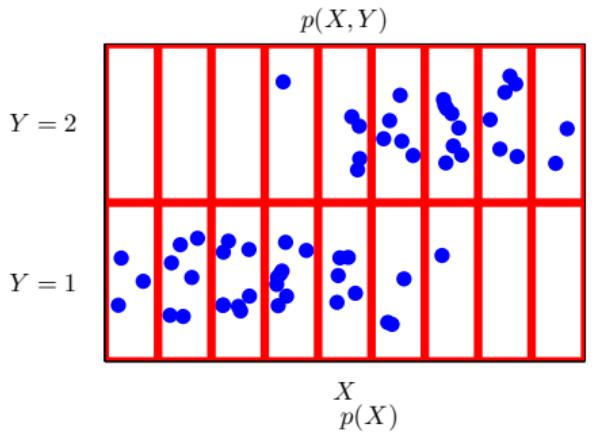
We can summarize the sum and product rules with a more compact notation:

$$\text{Sum rule: } p(X) = \sum_Y p(X, Y)$$

$$\text{Product rule: } p(X, Y) = p(Y|X)p(X)$$

These two rules underlie all the probabilistic machinery we'll use in this course.

Visualizing Joint Distributions



Bayes' Theorem

Using the symmetry $p(x, y) = p(Y, X)$ and the product rule:

$$\begin{aligned} p(X, Y) &= p(Y, X) \\ p(Y|X)p(X) &= p(X|Y)p(Y) \\ p(Y|X) &= \frac{p(X|Y)p(Y)}{p(X)} \end{aligned}$$

This is called *Bayes' Theorem* or *Bayes' Rule*.

Analysis of Medical Screening Example

With our probabilistic machinery we can now analyze our cancer screening example.

$$p(C = 1) = \frac{1}{100}$$

$$p(C = 0) = \frac{99}{100}$$

$$p(T = 1|C = 1) = \frac{90}{100}$$

$$p(T = 0|C = 1) = \frac{10}{100}$$

$$p(T = 1|C = 0) = \frac{3}{100}$$

$$p(T = 0|C = 0) = \frac{97}{100}$$

If we screen someone, probability that they test positive:

$$\begin{aligned} p(T = 1) &= p(T = 1|C = 0)p(C = 0) + p(T = 1|C = 1)p(C = 1) \\ &= \frac{3}{100} \times \frac{99}{100} + \frac{90}{100} \times \frac{1}{100} \\ &= \frac{387}{10,000} \\ &= .0387 \end{aligned}$$

If someone tests positive, probability they have cancer:

$$\begin{aligned} p(C = 1|T = 1) &= \frac{p(T = 1|C = 1)p(C = 1)p(C = 1)}{p(T = 1)} \\ &= \frac{90}{100} \times \frac{1}{100} \times \frac{10,000}{387} \\ &= \frac{90}{387} \\ &\approx 0.23 \end{aligned}$$

Prior and Posterior Probabilities

Bayes' Theorem updates our belief about someone's cancer.

- ▶ Before we run the test, the *prior probability* that someone has cancer is $p(C)$
- ▶ After we run the test, we use Bayes' Theorem to calculate the *posterior probability* $p(C|T)$

Independent Variables

If the joint distribution factorizes into the product of the marginals:

$$p(X, Y) = p(X)p(Y)$$

Then we say that X and Y are *independent*. So

$$P(Y|X) = p(Y)$$

and

$$P(X|Y) = p(X)$$

Question: in our cancer screening example, is the probability of a positive test independent of whether a person has cancer?