

# Reinforcement Learning

## Temporal-Difference Learning (RLbook 6)

Christopher Simpkins

Kennesaw State University




# Temporal-Difference Learning

- ▶ DP Review
  - ▶ Dynamic programming
  - ▶ Generalized policy iteration (GPI)
- ▶ Model-free control
  - ▶ Monte Carlo control
  - ▶ Temporal-difference learning

# Dynamic Programming


Dynamic programming algorithms, as well as RL algorithms in general, contain two phases (combined in value iteration):

- ▶ Prediction: estimate the value function
- ▶ Control: computing or approximating optimal policies



## Generalized Policy Iteration



In policy iteration we sweep entire state space in each step.



In generalized policy iteration (GPI) we interleave prediction and control at arbitrary granularity.

- ▶ As long as we visit every state, still assured of convergence.

# Monte Carlo Control



## Temporal-Difference Learning

Combination of dynamic programming and Monte Carlo ideas.

- ▶ Like Monte Carlo, learn directly from experience without a model of the environment.
- ▶ Like dynamic programming, update estimates based in part on other learned estimates, without waiting for a final outcome (bootstrap).

## TD Prediction

Use experience following a policy to update estimate of  $V$ , namely  $v_{pi}$ .

Monte Carlo methods wait until the return following the visit is known, then use that return as a target for  $V(S_t)$ :

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)] \quad (6.1)$$

TD methods only until the next time step. At time  $t + 1$  they immediately form a target and make a useful update using the observed reward  $R_{t+1}$  and the estimate  $V(S_{t+1})$

A simple TD method makes the update:

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + V(S_{t+1}) - \gamma V(S_t)]$$

immediate on transition to  $S_{t+1}$  and receiving  $R_{t+1}$ .

- ▶ Target for Monte Carlo update is  $G_t$ .
- ▶ Target for TD update is  $R_{t+1} + \gamma V(S_{t+1})$

# Tabular TD(0) Algorithm

One-step TD is called  $TD(0)$ , which is a special case of  $TD(\lambda)$  and  $n$ -step methods.

## Tabular TD(0) for estimating $v_\pi$

Input: the policy  $\pi$  to be evaluated

Algorithm parameter: step size  $\alpha \in (0, 1]$

Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

$A \leftarrow$  action given by  $\pi$  for  $S$

        Take action  $A$ , observe  $R, S'$

$V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$

$S \leftarrow S'$

    until  $S$  is terminal


## TD Error

Recall TD update:


$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + V(S_{t+1}) - \gamma V(S_t)]$$

The quantity in brackets is called *TD error* – the difference in the estimate value of  $S$  at time  $t$  and  $t + 1$ :

$$\delta \doteq R_{t+1} + V(S_{t+1}) - \gamma V(S_t)$$



## Sarsa: On-policy TD Control



Sarsa update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$



# Sarsa Algorithm

Sarsa (on-policy TD control) for estimating  $Q \approx q_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+$ ,  $a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

    Loop for each step of episode:

        Take action  $A$ , observe  $R, S'$


        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$$

$S \leftarrow S'; A \leftarrow A'$ ;

    until  $S$  is terminal

## Example: Windy Grid World



## Q-Learning: Off-policy TD Control

Sarsa update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(S_t, A_t)]$$

Q-learning update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha \left[ R_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(S_t, A_t) \right]$$

Q-learning is off-policy because the value update is made using  $\max_a$  rather than the  $a$  recommended by the policy being followed.

# Q-Learning Algorithm

Q-learning (off-policy TD control) for estimating  $\pi \approx \pi_*$

Algorithm parameters: step size  $\alpha \in (0, 1]$ , small  $\varepsilon > 0$

Initialize  $Q(s, a)$ , for all  $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$ , arbitrarily except that  $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

    Initialize  $S$

    Loop for each step of episode:

        Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\varepsilon$ -greedy)

        Take action  $A$ , observe  $R, S'$

$$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$$

$$S \leftarrow S'$$

    until  $S$  is terminal