

Introduction to Spark

Apache Spark

“Unified computing engine”

- ▶ cluster manager for running programs distributed across multiple computers (“nodes”)
 - ▶ Spark cluster manager, YARN, or Mesos
- ▶ libraries for parallel data processing

Spark Applications

- ▶ Driver process runs `main()` function on a node in the cluster
 - ▶ Maintains information about the Spark application
 - ▶ Responds to user's program or input
 - ▶ Analyzes, distributes, and schedules work across executors
- ▶ Executor processes
 - ▶ Carry out work assigned by driver
 - ▶ Report state of computation to driver




Figure 1: Spark Application Architecture

Spark Languages and APIs

Executors run Spark code. Driver runs user's program, which can be written in one of Spark's language APIs:

- ▶ Scala – the default language; Spark is written in Scala.
- ▶ Java
- ▶ Python – includes most of the Scala API
- ▶ SQL – subset of ANSI SQL 2003
- ▶ R – SparkR is part of Spark core, R community provides an alternative called spraklyr




Figure 2: Spark Session and Language APIs

Note: there are two APIs: a high-level “structured” API, and a low-level API.

Spark Session

Every Spark application has exactly one associated `SparkSession`.

- ▶ Explicitly create a `SparkSession` in Spark programs
- ▶ `SparkSession` instance implicitly created in Spark shell

```
1 scala> spark
2 res0: org.apache.spark.sql.SparkSession =
    org.apache.spark.sql.SparkSession@32950acc
```


DataFrames

A DataFrame is table of rows and columns partitioned across the nodes in the Spark cluster.

Spreadsheet on
a single machine



Table or Data Frame
partitioned across servers
in a data center



DataFrame

```
1 scala> val myRange = spark.range(1000).toDF("number")
2 myRange: org.apache.spark.sql.DataFrame = [number: bigint]
3
4 scala> myRange.show(5)
5 +-----+
6 |number|
7 +-----+
8 |    0|
9 |    1|
10 |    2|
11 |    3|
12 |    4|
13 +-----+
14 only showing top 5 rows
```

A DataFrame's schema is a list of columns and their types:

```
1 scala> myRange.schema
2 res6: org.apache.spark.sql.types.StructType =
      StructType(StructField(number,LongType,false))
3
4 scala> myRange.printSchema
5 root
6 |-- number: long (nullable = false)
```

A DataFrame is

DataFrame Computation Pipeline

- ▶ Transformations take a DataFrame as input and produce another DataFrame (lazily)
- ▶ An action is executed as the last step of a computation pipeline to produce the final DataFrame