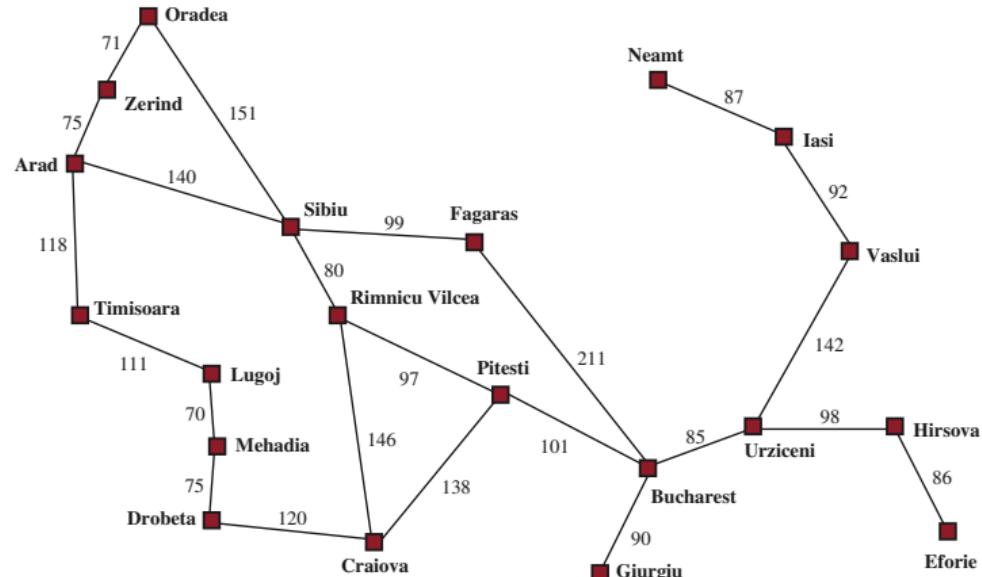


Problem Solving

Artificial Intelligence

Christopher Simpkins

Problem-Solving Agents



- ▶ In this lesson we consider a *state* to be our location in one of these cities.
- ▶ A *goal* is a state in which we are located in a particular city.

This is the essence of problem solving: transforming a current state into a goal state. The first family of algorithms we'll study for problem solving are *search* algorithms.

Problem Solving Process

To solve a problem, we

- ▶ Formulate a **goal**, e.g., “reach Bucharest”
- ▶ Formulate the **problem** as a set of states and actions that move us from one state to another.
 - ▶ Problem is a **model** – an *abstract* mathematical description.
 - ▶ Abstraction is essence and ignorance.
 - ▶ Key skill in problem formulation is finding the right **level of abstraction**.
- ▶ **Search** the possible sequences of action in our problem model that transforms our state from the current state to the goal state. A sequence of actions that gets us to the goal state is called a *solution*. May be many; pick one.
- ▶ **Execute** the actions in the solution.

Open-Loop vs. Closed-Loop

- ▶ In an **open-loop** system the agent gets no feedback, i.e., sensor input, after executing an action.
 - ▶ If the agent's model is perfect and actions are deterministic, then the agent can operate in an open-loop fashion, simply executing the actions in the solution one after the other.
- ▶ In a **closed-loop** system gets sensory feedback after every action, so it can check whether the action had the expected effect.
 - ▶ If the environment is partially observable or actions are nondeterministic, closed-loop control is necessary.
 - ▶ Say the agent executes to `ToSibiu` action but ends up in `Zerind`. Closed-loop feedback will alert the agent to this fact so it can re-plan.

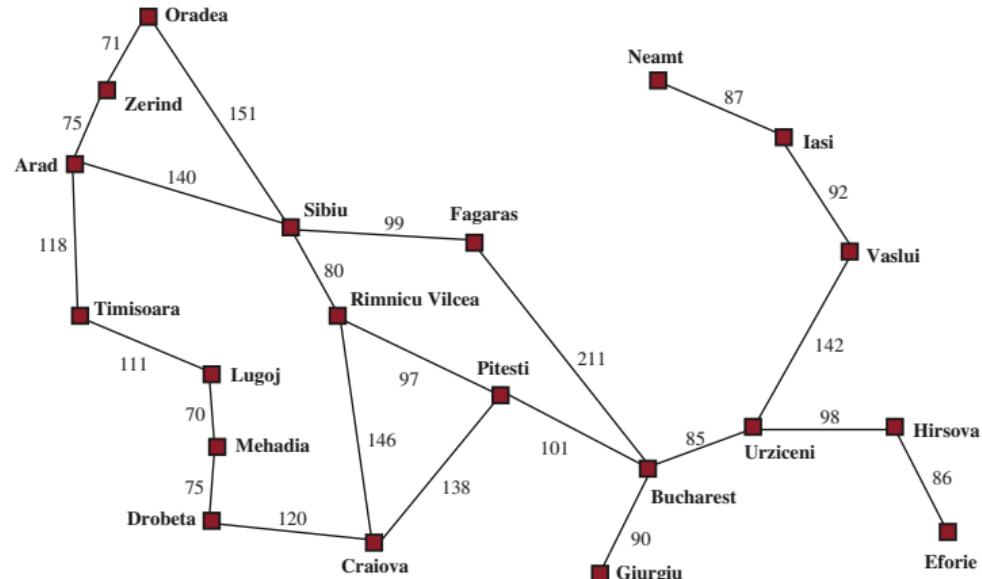
Search Problems and Solutions

A search problem consists of:

- ▶ A set of **states**, which we call a **state space**.
- ▶ **Initial state**
- ▶ A set of **goal states**.
 - ▶ Typically use an **IS-GOAL(s)** predicate function to identify goal states.
- ▶ Sets of **actions** available in each state, **ACTION(s)**
 - ▶ **ACTION(Arad)= {ToSibiu, ToTimisoara, ToZerind}**
- ▶ A **transition model**, **RESULT(s, a)**
 - ▶ **RESULT(Arad, ToZerind)= Zerind**
- ▶ An **action cost function**, **ACTION-COST(s, a, s')** or $c(s, a, s')$ which returns the cost of executing action a in state s and reaching state s' .

Solution

- ▶ A solution is a path from the start state to the a goal state.
- ▶ An optimal solution is a solution with lowest cost among all solutions.




- ▶ How many paths are there from **Arad** to **Bucharest**?
- ▶ What is/are the solutions to the **Arad-to-Bucharest** problem (assume perfect information – fully observable, known dynamics, and deterministic actions)?

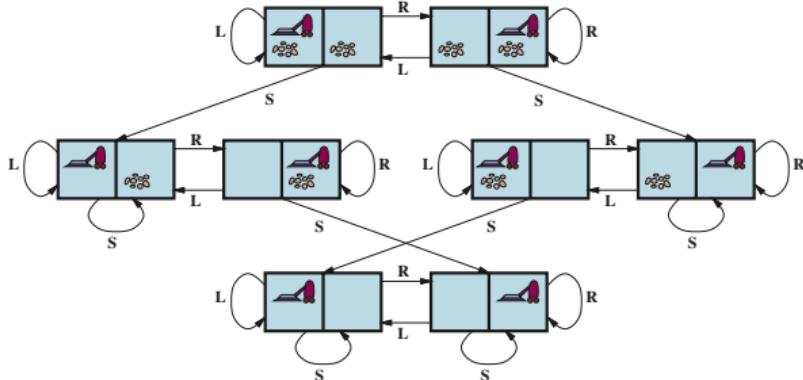
Example Problems

- ▶ **Standardized problems** use idealized environments designed to illustrate or exercise various problem-solving methods. See, for example, [Gymnasium](#).
 - ▶ A **grid world** is an standardized environment whose states are organized as a grid, and whose actions include moving between adjacent grids.
- ▶ **Real-world problems** are formulated for specific real-world tasks, like the problem specification used for Rhoombas.

Here's a standardized environment for the vacuum cleaner agent, formulated as a grid world:



Vacuum Cleaner Grid World



- ▶ **States** include both the agent's location, and characteristics of the environment. For the vacuum world, that's $2 \cdot 2^2 = 8$ states.
- ▶ **Initial state** is an arbitrary choice of the possible states. Sometimes this choice is important.
- ▶ **Actions** for this vacuum world are are **L**, **R**, and **Suck**.
 - ▶ For 2D grids we can choose between
 - ▶ **absolute movement**, like **Up** and **Right**, a.k.a., cardinal directions, or
 - ▶ **egocentric movement**, like **TurnRight**, **MoveForward**. How does this affect the state description?
- ▶ **Goal states** are those in which every location is clean.
- ▶ **Action cost** (path cost) is 1.

Route Finding

- ▶ **States:** a location (e.g., an airport) and the time.
 - ▶ If action cost (e.g., a flight segment) depends on previous segments, fares, etc., the state must include these details.
- ▶ **Initial state:** The user's home airport.
- ▶ **Actions:** Take any flight from the current location, in any seat class, leaving after the current time, or for connecting flights, after sufficient in-airport transfer time.
- ▶ **Transition model:** The state resulting from taking a flight will have the flight's destination as the new location and the flight's arrival time as the new time.
 - ▶ Example $T(s, a, s')$: $T(\mathbf{S(ATL, 10:00)}, \mathbf{A(DL875)}, \mathbf{S(LGA, 12:00)})$ (DL875 has a flight time of 2 hours).
- ▶ **Goal state:** A destination city. Sometimes the goal can be more complex, such as arrive at the destination on a nonstop flight. (Remember, a solution is a path, i.e., sequence of actions.)
- ▶ **Action cost:** A combination of monetary cost, waiting time, flight time, customs and immigration procedures, seat quality, time of day, type of airplane, frequent-flyer reward points, and so on.

Real-World Problems

- ▶ **Touring problems**
- ▶ **VLSI layout** – minimize area, minimize circuit delays, minimize stray capacitances, and maximize manufacturing yield
 - ▶ Cell layout – place cells on chip so they don't overlap and have room for connections
 - ▶ Channel routing – find routes for each wire between cells
- ▶ **Robot navigation**
- ▶ **Automatic assembly sequencing** – standard practice in manufacturing since the 1970s.
 - ▶ Solving some automatic assembly problems could earn you a [Nobel Prize!](#)

Search Algorithms


A **search algorithm** takes a search problem as input and returns a solution, or an indication of failure.

- ▶ In general, the states and actions of a problem create a state space graph.
- ▶ Here we consider algorithms that superimpose a **search tree** over the state-space graph.
- ▶ **Nodes** correspond to states, **edges** correspond to actions

Don't confuse state space with search tree.

- ▶ State space is set of states, and actions that cause transitions between states.
- ▶ Search tree describes paths between these states, reaching towards the goal(s).
 - ▶ May be many nodes for a given state, but each path from root to node is unique.

Here is a search tree being imposed on the Romania state space graph by a search algorithm.



Elements of Search Algorithms


Essence of search:

- ▶ Choose a child node to consider next. “Who’s first?”
- ▶ Put aside other nodes for later. “Who’s next?”

Root node is initial state. At each node we can **expand** the node, which grows the tree, by taking actions (adding edges) that lead to successor states (generate successor/child nodes). Search algorithms must keep track of:

- ▶ *Expanded* nodes. We test expanded nodes before dealing with frontier.
- ▶ *Frontier* nodes, which are generated but not yet expanded.
 - ▶ $\text{Reached} = \text{Expanded} + \text{Frontier}$


Search Progression




- ▶ Expanded nodes are lavender with bold letters.
- ▶ Frontier nodes are green with normal weight font.
- ▶ Nodes in dashed-line ovals are candidates for expansion.

Separation Property of Graph Search


The **frontier** separates the interior region of expanded nodes from the exterior region of unexpanded nodes.



(a)



(b)



(c)

Frontier is in green. Interior is lavender. Exterior is faint dashed.

- ▶ (a) Only root expanded.
- ▶ (b) Top frontier node expanded.
- ▶ (c) Remaining successors of root expanded in clockwise order.

Implementation Note: The `yield` statement

A function containing a `yield` statement is a **generator**. Use a generator to turn a data generating process into an iterator. Node expansion is a data generating process.

```
1 In [36]: def by_twos(start: int, end: int):
2     ...:     x = start
3     ...:     while x < end:
4     ...:         yield x
5     ...:         x += 2
6     ...:
7
8 In [37]: by_twos(1, 9)
9 Out[37]: <generator object by_twos at 0x109010ee0>
10
11 In [38]: list(Out[37])
12 Out[38]: [1, 3, 5, 7]
13
14 In [39]: for x in by_twos(1, 10):
15     ...:     print(f"{x=}")
16     ...:
17 x=1
18 x=3
19 x=5
20 x=7
21 x=9
```

Search Data Structures

Node:

- ▶ `node.STATE`: the state to which the node corresponds;
- ▶ `node.PARENT`: the node in the tree that generated this node;
- ▶ `node.ACTION`: the action that was applied to the parent's state to generate this node;
- ▶ `node.PATH-COST`: the total cost of the path from the initial state to this node. In mathematical formulas, we use $g(node)$ as a synonym for PATH-COST.

Frontier is a **queue** with operations:

- ▶ `IS-EMPTY(frontier)` returns true only if there are no nodes in the frontier.
- ▶ `POP(frontier)` removes the top node from the frontier and returns it.
- ▶ `TOP(frontier)` returns (but does not remove) the top node of the frontier.
- ▶ `ADD(node, frontier)` inserts node into its proper place in the queue.

Queues used in search algorithms:

- ▶ A **priority queue** first pops the node with the minimum cost according to some evaluation function, f . It is used in best-first search.
- ▶ A **FIFO queue** or first-in-first-out queue first pops the node that was added to the queue first; we shall see it is used in breadth-first search.
- ▶ A **LIFO queue** or last-in-first-out queue (also known as a stack) pops first the most recently added node; we shall see it is used in depth-first search.

Best-First Search

Best-first search is an abstract search algorithm. Name can be tricky to understand.

- ▶ Best way to pick the *first* node to consider next.
- ▶ We use a generalization of queues, called a *priority queue*, to store the *frontier*.
- ▶ An evaluation function, $f(node)$, imposes an ordering on the nodes in the priority queue.

The evaluation function considers the path to the node, not any property of the node itself. Remember, a solution to a search problem is characterized by the path from the root to the goal, not some characteristic of the goal.


We'll now describe several uninformed search algorithms. I recommend you also look at their [implementations in Python](#), which may be easier to follow.

Best-First Search Algorithm

```
function BEST-FIRST-SEARCH(problem,f) returns a solution node or failure
    node  $\leftarrow$  NODE(STATE=problem.INITIAL)
    frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
    reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
    while not Is-EMPTY(frontier) do
        node  $\leftarrow$  POP(frontier)
        if problem.IS-GOAL(node.STATE) then return node
        for each child in EXPAND(problem, node) do
            s  $\leftarrow$  child.STATE
            if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
                reached[s]  $\leftarrow$  child
                add child to frontier
    return failure
```

```
function EXPAND(problem, node) yields nodes
    s  $\leftarrow$  node.STATE
    for each action in problem.ACTIONS(s) do
        s'  $\leftarrow$  problem.RESULT(s, action)
        cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
        yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```

Redundant Paths



In the path from **Arad** to **Sibiu** to **Arad**,

- ▶ **Arad** is a **repeated state** and
- ▶ the path is a **cycle, or loopy path**.

Cycle special case of **redundant path**: multiple paths to the same state. Three approaches:

1. Remember reaches states, like best-first search. Best when reached states fits in memory.
2. Don't worry about repeated states. Works when repeated states rare or impossible.
 - ▶ **Graph search** checks for redundant paths, which occur in graphs in general.
 - ▶ **Tree-like search** does not check for redundant paths, since trees are acyclic graphs.
3. Only check for cycles, not other kinds of redundant paths.
 - ▶ E.g., search path in reverse

Measuring Problem-Solving Performance

- ▶ **Completeness:** Is the algorithm guaranteed to find a solution when there is one, and to correctly report failure when there is not?
 - ▶ Complete search algorithms must be **systematic**.
 - ▶ Easier to achieve for finite state spaces.
 - ▶ In an infinite state space with no solution, search won't terminate.
- ▶ **Cost optimality:** Does it find a solution with the lowest path cost of all solutions?
- ▶ **Time complexity:** How long does it take to find a solution? This can be measured in seconds, or more abstractly by the number of states and actions considered.
- ▶ **Space complexity:** How much memory is needed to perform the search?

For *explicit* graphs, like Romania, time and space complexity typically expressed in terms of number of vertices (state nodes), $|V|$, and number of edges, $|E|$ (state-action pairs, which generate $((s, a, s'))$ triples).

For *implicit* state space graphs we characterize time and space complexity in terms of depth, d (number of actions in an optimal solution), and branching factor, b (number of successor nodes per node).

Uninformed Search Strategies

Uninformed search strategies have no information about which actions are better for reaching a goal. In these cases we can only do systematic searches of the state space. We'll discuss

- ▶ Breadth-first search
- ▶ Uniform-Cost search (Dijkstra's algorithm)
- ▶ Depth-first search
- ▶ Depth-limited searchand
- ▶ Iterative deepening search.
- ▶ Bidirectional search

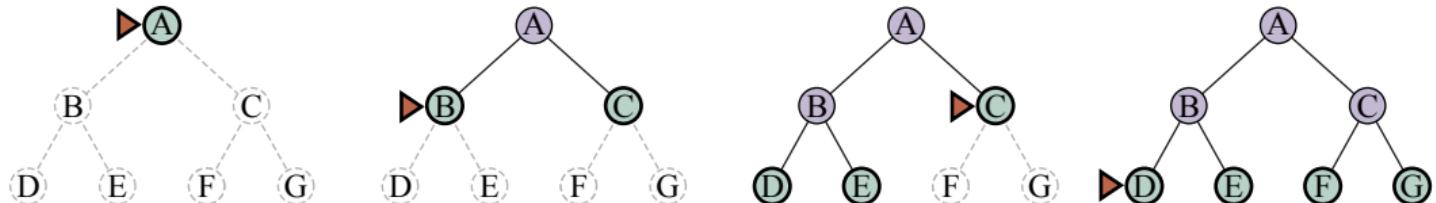
Breadth-First Search

- ▶ Good when path costs are uniform.
- ▶ Equivalent to best-first search where $f(\text{node})$ is the depth of the node
- ▶ Guaranteed to find minimal number of actions because it evaluates depth d before generating depth $d + 1$.

But three optimizations afforded by the BFS algorithm and uniform path costs:

- ▶ FIFO queue instead of priority queue
- ▶ *Reached* is a set instead of a mapping $S \rightarrow \text{Node}$
 - ▶ With uniform path costs, as soon as BFS finds a node, it's the fastest way to it.
- ▶ **Early goal test** – as soon as we expand a node, we can test it.

BFS Algorithm



```
function BREADTH-FIRST-SEARCH(problem) returns a solution node or failure
    node  $\leftarrow$  NODE(problem.INITIAL)
    if problem.IS-GOAL(node.STATE) then return node
    frontier  $\leftarrow$  a FIFO queue, with node as an element
    reached  $\leftarrow \{\text{problem.INITIAL}\}
    while not IS-EMPTY(frontier) do
        node  $\leftarrow$  POP(frontier)
        for each child in EXPAND(problem, node) do
            s  $\leftarrow$  child.STATE
            if problem.IS-GOAL(s) then return child
            if s is not in reached then
                add s to reached
                add child to frontier
    return failure$ 
```

Analysis of BFS


- ▶ Complete, because it generates all nodes at each depth.
- ▶ Time complexity: at each level, b nodes for each b predecessors, so
 - ▶ $1 + b + b^2 + b^3 + \dots + b^d = O(b^d)$
- ▶ Space complexity: $O(b^d)$ because all nodes are stored while the search proceeds.

Uninformed search is not appropriate for exponential complexity problems except for smallest instances. Assuming your computer can process 1 million nodes per second and store each node in 1 Kb,

- ▶ For a problem with $b = 10$ and $d = 10$, how long will it take search and how much space will be required?
- ▶ Same problem, but with $d = 14$?

Uniform-Cost Search (Dijkstra's Algorithm)

BFS where the best-first $f(node)$ is the path cost to the current node.



```
function BEST-FIRST-SEARCH(problem, f) returns a solution node or failure
    node  $\leftarrow$  NODE(STATE=problem.INITIAL)
    frontier  $\leftarrow$  a priority queue ordered by f, with node as an element
    reached  $\leftarrow$  a lookup table, with one entry with key problem.INITIAL and value node
    while not IS-EMPTY(frontier) do
        node  $\leftarrow$  POP(frontier)
        if problem.IS-GOAL(node.STATE) then return node
        for each child in EXPAND(problem, node) do
            s  $\leftarrow$  child.STATE
            if s is not in reached or child.PATH-COST < reached[s].PATH-COST then
                reached[s]  $\leftarrow$  child
                add child to frontier
    return failure


function EXPAND(problem, node) yields nodes
    s  $\leftarrow$  node.STATE
    for each action in problem.ACTIONS(s) do
        s'  $\leftarrow$  problem.RESULT(s, action)
        cost  $\leftarrow$  node.PATH-COST + problem.ACTION-COST(s, action, s')
        yield NODE(STATE=s', PARENT=node, ACTION=action, PATH-COST=cost)
```

Analysis of Uniform-Cost Search

Let C^* be the cost of the optimal solution and $\epsilon > 0$ be a lower bound on the cost of each action.

- ▶ Time and space complexity are $O(b^{1+\lfloor \frac{C^*}{\epsilon} \rfloor})$.
 - ▶ Since lower cost paths are always explored first, even when a higher cost path might be the one to lead to an optimal solution, can be worse than BFS.
 - ▶ If all action costs equal, then it's like BFS, $O(b^{1+d})$.
- ▶ Complete. like BFS
- ▶ Cost-optimal, because a solution will be at least as low cost as any other in the frontier.

Depth-First Search




Depth-Limited Search and Iterative Deepening Search

```
function ITERATIVE-DEEPENING-SEARCH(problem) returns a solution node or failure
  for depth = 0 to  $\infty$  do
    result  $\leftarrow$  DEPTH-LIMITED-SEARCH(problem, depth)
    if result  $\neq$  cutoff then return result

function DEPTH-LIMITED-SEARCH(problem,  $\ell$ ) returns a node or failure or cutoff
  frontier  $\leftarrow$  a LIFO queue (stack) with NODE(problem.INITIAL) as an element
  result  $\leftarrow$  failure
  while not IS-EMPTY(frontier) do
    node  $\leftarrow$  POP(frontier)
    if problem.IS-GOAL(node.STATE) then return node
    if DEPTH(node)  $>$   $\ell$  then
      result  $\leftarrow$  cutoff
    else if not IS-CYCLE(node) do
      for each child in EXPAND(problem, node) do
        add child to frontier
  return result
```

Progression of Iterative Deepening Search



Bidirectional Best-First Search

```
function BiBF-SEARCH(problemF, fF, problemB, fB) returns a solution node, or failure
    nodeF  $\leftarrow$  NODE(problemF.INITIAL) // Node for a start state
    nodeB  $\leftarrow$  NODE(problemB.INITIAL) // Node for a goal state
    frontierF  $\leftarrow$  a priority queue ordered by fF, with nodeF as an element
    frontierB  $\leftarrow$  a priority queue ordered by fB, with nodeB as an element
    reachedF  $\leftarrow$  a lookup table, with one key nodeF.STATE and value nodeF
    reachedB  $\leftarrow$  a lookup table, with one key nodeB.STATE and value nodeB
    solution  $\leftarrow$  failure
    while not TERMINATED(solution, frontierF, frontierB) do
        if fF(TOP(frontierF)) < fB(TOP(frontierB)) then
            solution  $\leftarrow$  PROCEED(F, problemF, frontierF, reachedF, reachedB, solution)
        else solution  $\leftarrow$  PROCEED(B, problemB, frontierB, reachedB, reachedF, solution)
    return solution

function PROCEED(dir, problem, frontier, reached, reached2, solution) returns a solution
    // Expand node on frontier; check against the other frontier in reached2.
    // The variable "dir" is the direction: either F for forward or B for backward.
    node  $\leftarrow$  POP(frontier)
    for each child in EXPAND(problem, node) do
        s  $\leftarrow$  child.STATE
        if s not in reached or PATH-COST(child) < PATH-COST(reached[s]) then
            reached[s]  $\leftarrow$  child
            add child to frontier
            if s is in reached2 then
                solution2  $\leftarrow$  JOIN-NODES(dir, child, reached2[s]))
                if PATH-COST(solution2) < PATH-COST(solution) then
                    solution  $\leftarrow$  solution2
    return solution
```

Comparing Uninformed Search Algorithms

Criterion	Breadth-First	Uniform-Cost	Depth-First	Depth-Limited	Iterative Deepening	Bidirectional (if applicable)
Complete?	Yes ¹	Yes ^{1,2}	No	No	Yes ¹	Yes ^{1,4}
Optimal cost?	Yes ³	Yes	No	No	Yes ³	Yes ^{3,4}
Time	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(b^m)$	$O(b^\ell)$	$O(b^d)$	$O(b^{d/2})$
Space	$O(b^d)$	$O(b^{1+\lfloor C^*/\epsilon \rfloor})$	$O(bm)$	$O(b\ell)$	$O(bd)$	$O(b^{d/2})$