

Reinforcement Learning

Temporal-Difference Learning (RLbook 6)

Christopher Simpkins

Kennesaw State University


1 Temporal-Difference Learning

- DP Review
 - Dynamic programming
 - Generalized policy iteration (GPI)
- Model-free control
 - Monte Carlo control
 - Temporal-difference learning

2 Dynamic Programming


Dynamic programming algorithms, as well as RL algorithms in general, contain two phases (combined in value iteration):

- Prediction: estimate the value function
- Control: computing or approximating optimal policies



3 Generalized Policy Iteration


In policy iteration we sweep entire state space in each step.



In generalized policy iteration (GPI) we interleave prediction and control at arbitrary granu-

larity.

- As long as we visit every state, still assured of convergence.



4 Monte Carlo Methods

Monte Carlo methods sample and average returns for each state–action pair.

- Generate a set of episode samples by following a policy π
- For each visit to a state s , find the first visit to s , then average all the returns following the first visit to s to estimate the value of s

Combination of dynamic programming and Monte Carlo ideas.

- Like Monte Carlo, learn directly from experience without a model of the environment.
- Like dynamic programming, update estimates based in part on other learned estimates, without waiting for a final outcome (bootstrap).

5 TD Prediction

Use experience following a policy to update estimate of V , namely v_{pi} .

Monte Carlo methods wait until the return following the visit is known, then use that return as a target for $V(S_t)$:

$$V(S_t) \leftarrow V(S_t) + \alpha [G_t - V(S_t)] \quad (6.1)$$

TD methods only until the next time step. At time $t + 1$ they immediately form a target and make a useful update using the observed reward R_{t+1} and the estimate $V(S_{t+1})$

A simple TD method makes the update:

$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

immediate on transition to S_{t+1} and receiving R_{t+1} .

- Target for Monte Carlo update is G_t .
- Target for TD update is $R_{t+1} + \gamma V(S_{t+1})$

6 Tabular TD(0) Algorithm

One-step TD is called $TD(0)$, which is a special case of $TD(\lambda)$ and n -step methods.

Tabular TD(0) for estimating v_π

```
Input: the policy  $\pi$  to be evaluated
Algorithm parameter: step size  $\alpha \in (0, 1]$ 
Initialize  $V(s)$ , for all  $s \in \mathcal{S}^+$ , arbitrarily except that  $V(\text{terminal}) = 0$ 
Loop for each episode:
  Initialize  $S$ 
  Loop for each step of episode:
     $A \leftarrow$  action given by  $\pi$  for  $S$ 
    Take action  $A$ , observe  $R, S'$ 
     $V(S) \leftarrow V(S) + \alpha [R + \gamma V(S') - V(S)]$ 
     $S \leftarrow S'$ 
  until  $S$  is terminal
```


7 TD Error

Recall TD update:


$$V(S_t) \leftarrow V(S_t) + \alpha [R_{t+1} + \gamma V(S_{t+1}) - V(S_t)]$$

The quantity in brackets is called **TD error** – the difference in the estimate value of S at time t and $t + 1$:

$$\delta \doteq R_{t+1} + \gamma V(S_{t+1}) - V(S_t)$$




8 Sarsa: On-policy TD Control



Sarsa update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$



9 Sarsa Algorithm

Sarsa (on-policy TD control) for estimating $Q \approx q_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

Initialize S

Choose A from S using policy derived from Q (e.g., ε -greedy)

Loop for each step of episode:

Take action A , observe R, S'


Choose A' from S' using policy derived from Q (e.g., ε -greedy)

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma Q(S', A') - Q(S, A)]$

$S \leftarrow S'; A \leftarrow A';$

until S is terminal

10 Example: Windy Grid World



11 Q-Learning: Off-policy TD Control

Sarsa update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

Q-learning update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha [R_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

Q-learning is off-policy because the value update is made using $\max_a$ rather than the a recommended by the policy being followed.

12 Q-Learning Algorithm

Q-learning (off-policy TD control) for estimating $\pi \approx \pi_*$

Algorithm parameters: step size $\alpha \in (0, 1]$, small $\varepsilon > 0$

Initialize $Q(s, a)$, for all $s \in \mathcal{S}^+, a \in \mathcal{A}(s)$, arbitrarily except that $Q(\text{terminal}, \cdot) = 0$

Loop for each episode:

 Initialize S

 Loop for each step of episode:

 Choose A from S using policy derived from Q (e.g., ε -greedy)

 Take action A , observe R, S'

$Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$

$S \leftarrow S'$

 until S is terminal