

Artificial Intelligence

Uncertainty

Christopher Simpkins

Kennesaw State University

Acting Under Uncertainty

Logical agents maintain a belief state and generate contingency plans that account for every possibility. This breaks down for nontrivial problems due to:

- ▶ Exhaustivity of belief state, which must contain all possible states, even unlikely states.
- ▶ Exhaustivity of contingency plan, which must account for every possible, however unlikely, action outcome.
- ▶ Unsatisfiability. There may be no guaranteed plan to achieve the goal. If we must act anyway, how do we choose the best plan?
- ▶ Qualification problem. The closed-world assumption allows us to simplify logical environment specifications for simple domains, but real-world domains contain far more detail which must be accounted for.

Rational decision under uncertainty depends on:

- ▶ relative importance of goals,
- ▶ likelihood of achieving goals, and
- ▶ degree to which goals can be achieved.

From Logic to Probability Theory

Consider:

Toothache \Rightarrow Cavity

Not all patients with toothaches have cavities. How about:

Toothache \Rightarrow Cavity \vee GumProblem \vee Abscess ...

We'd need large list after the dots. Try causal direction:

Cavity \Rightarrow Toothache

But not all cavities hurt. Logic fails to deal with complex domains like medical diagnosis due to:

- ▶ **Exhaustivity** (laziness). Too many antecedents or consequents to list.
- ▶ **Theoretical ignorance**. Rare to have complete logical theory of nontrivial domain.
- ▶ **Practical ignorance**. Even with a complete domain theory, hard to measure all the necessary inputs (medical tests, etc.)

We need to deal with uncertain knowledge, with **degrees of belief**. For that, we use probability theory.

Probability Theory as Knowledge Representation

- ▶ Ontological commitments of logic and probability theory same:
 - ▶ World composed of facts that do or do not hold in any particular case.
- ▶ Epistemological commitments differ:
 - ▶ Logic assigns true, false, or no opinion to each sentence.
 - ▶ Probability theory assigns numerical degree of belief between 0 (for sentences that are certainly false) and 1 (certainly true) to each sentence.

Probability theory solves the qualification problem by summarizing the uncertainty stemming from our laziness and ignorance.

What do probabilities mean?



Probabilities are about our uncertain knowledge.

"80% chance (probability 0.8) patient with a toothache has a cavity."

- ▶ Out of all the situations that are indistinguishable from the current situation *as far as our knowledge goes*, the patient will have a cavity in 80% of them.
- ▶ Belief could come from statistical data, or domain theory, or combination of sources.

There is no uncertainty in the actual world. The patient either has a cavity or does not. Probabilities refer to our knowledge of the world state, not the actual world state.

If our knowledge changes, e.g., we find out patient has history of gum disease, we make a different statement.

"Given patient has a history of gum disease and a toothache, 40% chance patient has gum disease."

Rational Decisions

How do we choose between different plans that achieve the goal?

- ▶ Each plan has a likelihood of success, e.g.:
 - ▶ Plan A has a 95% chance of achieving a goal state.
- ▶ Each plan may lead to goal states with different outcomes, each of which is a goal state.
 - ▶ An **outcome** is a completely specified state.

Elements of outcomes are more or less desirable – more or less *useful* – to a given agent:

- ▶ utility: the quality of being useful.

Utility theory associates a utility value with each possible outcome, which induces a preference ordering over outcomes.

An agent is rational if and only if it chooses the action that yields the highest expected utility, averaged over all the possible outcomes of the action.

Decision-Theoretic Agents

Decision theory = probability theory + utility theory

```
function DT-AGENT(percept) returns an action
    persistent: belief-state, probabilistic beliefs about the current state of the world
                action, the agent's action

    update belief-state based on action and percept
    calculate outcome probabilities for actions,
        given action descriptions and current belief-state
    select action with highest expected utility
        given probabilities of outcomes and utility information
    return action
```

Each action, a , leads to a probability distribution over outcomes, or “result states”:

$$\sum_{s \in S} Pr(s) = 1$$

And each outcome state has a utility. A rational decision-theoretic agent chooses the action that maximizes expected utility, that is:

$$\operatorname{argmax}_a \sum_{s \in S} Pr(s)U(s)$$

Probability Models

In logic we have satisfiability: $M(\alpha)$ is the set of all possible worlds in which sentence α is true.

In probability theory the set of possible worlds is a **sample space**, Ω , which must be mutually exclusive and exhaustive (one and only one possible world must be the case).

$$0 \leq Pr(\omega) \leq 1 \text{ for every } \omega \text{ and } \sum_{\omega \in \Omega} Pr(\omega) = 1$$

An **event** is a set of possible worlds, e.g., “probability two dice sum to 6.”

A **proposition** is an event expressed in a formal language; specifically, for each proposition, the corresponding set contains just those possible worlds in which the proposition holds.

The probability associated with a proposition is defined to be the sum of the probabilities of the worlds in which it holds:

$$\text{For any proposition } \phi, Pr(\phi) = \sum_{\omega \in \phi} Pr(\omega)$$

Example: $Pr(Total = 11) = Pr((5, 6)) + Pr((6, 5)) = \frac{1}{36} + \frac{1}{36} = \frac{1}{18}$

Prior and Conditional Probabilities

An **unconditional** or **prior** probability is a degree of belief in a proposition with no other information.

A **conditional** or **posterior** probability is a degree of belief in a proposition given some relevant information.

Definition of conditional probability:

$$Pr(a|b) = \frac{Pr(a \wedge b)}{Pr(b)}$$

From that definition we get the **product rule**:

$$Pr(a \wedge b) = Pr(a|b)Pr(b)$$

Conditional probability is not implication.

The assertion $Pr(\text{cavity}|\text{toothache}) = 0.6$ does not mean “Whenever toothache is true, conclude that cavity is true with probability 0.6”

It means “Whenever toothache is true *and we have no further information*, conclude that cavity is true with probability 0.6.”

If we had the further information that the dentist found no cavities, certainly not the case that cavity is true with probability 0.6; instead we have $Pr(\text{cavity}|\text{toothache and } \neg\text{cavity}) = 0$.

Probability Assertions

A **random variable**, which begins with capital letter, is a function mapping from possible worlds Ω to a set of possible values the variable can take.

- ▶ $Die_1 : \{1, \dots, 6\} \rightarrow \{1, \dots, 6\}$
- ▶ $Doubles : \{1, \dots, 6\} \times \{1, \dots, 6\} \rightarrow \{\text{true}, \text{false}\}$
- ▶ $Weather : \{\text{sun}, \text{rain}, \text{cloud}, \text{snow}\} \rightarrow \{\text{sun}, \text{rain}, \text{cloud}, \text{snow}\}$

Individual values are written in lowercase and often abbreviated.

- ▶ $Pr(\text{sun})$ stands for $Pr(Weather = \text{sun})$.

A **probability distribution** is an assignment of probabilities to all values of a random variable,
e.g.:

$$Pr(Weather = \text{sun}) = 0.6$$

$$Pr(Weather = \text{rain}) = 0.1$$

$$Pr(Weather = \text{cloud}) = 0.29$$

$$Pr(Weather = \text{snow}) = 0.01$$

Joint Probability Distributions

$Pr(Weather, Cavity)$ denotes the probabilities of all combinations of *Weather* and *Cavity*, which is called a **joint probability distribution**.

Probability Axioms

All of probability theory can be built from **Kolmogorov's axioms**.

Law of total probability:

$$0 \leq Pr(\omega) \leq 1 \text{ for every } \omega \text{ and } \sum_{\omega \in \Omega} Pr(\omega) = 1$$

Probability of a disjunction, **inclusion-exclusion principle**:

$$Pr(a \vee b) = Pr(a) + Pr(b) - Pr(a \wedge b)$$

Why is probability theory a valid basis for rational behavior?

De Finetti's Theorem:

If Agent 1 expresses a set of degrees of belief that violate the axioms of probability theory then there is a combination of bets by Agent 2 that guarantees that Agent 1 will lose money every time.

Example:

Proposition	Agent 1's belief	Agent 2 bets	Agent 1 bets	Agent 1 payoffs for each outcome			
				a, b	$a, \neg b$	$\neg a, b$	$\neg a, \neg b$
a	0.4	\$4 on a	\$6 on $\neg a$	-\$6	-\$6	\$4	\$4
b	0.3	\$3 on b	\$7 on $\neg b$	-\$7	\$3	-\$7	\$3
$a \vee b$	0.8	\$2 on $\neg(a \vee b)$	\$8 on $a \vee b$	\$2	\$2	\$2	-\$8
				-\$11	-\$1	-\$1	-\$1

Inference Using Full Joint Distributions

Knowledge base is full joint distribution of boolean random variables *Toothache*, *Cavity*, *Catch*.

		<i>toothache</i>	\neg <i>toothache</i>		
		<i>catch</i>	\neg <i>catch</i>	<i>catch</i>	\neg <i>catch</i>
<i>cavity</i>		0.108	0.012	0.072	0.008
\neg <i>cavity</i>		0.016	0.064	0.144	0.576

- ▶ All probabilities above sum to 1.
- ▶ $Pr(cavity \vee toothache) = 0.108 + 0.012 + 0.072 + 0.008 + 0.016 + 0.064 = 0.28$.

Unconditional or **marginal probability** of cavity:

$$Pr(cavity) = 0.108 + 0.012 + 0.072 + 0.008 = 0.2$$

Marginalization

<i>toothache</i>		\neg <i>toothache</i>	
	<i>catch</i>	\neg <i>catch</i>	<i>catch</i>
<i>cavity</i>	0.108	0.012	0.072
\neg <i>cavity</i>	0.016	0.064	0.144
			0.576

General rule: $Pr(Y) = \sum_z Pr(Y, Z = z)$ where Z is set of all possible values of variables other than Y . Example, let $Y = Cavity$

$$\begin{aligned} Pr(Cavity) &= Pr(Cavity, toothache, catch) + Pr(Cavity, toothache, \neg catch) \\ &\quad + Pr(Cavity, \neg toothache, catch) + Pr(Cavity, \neg toothache, \neg catch) \\ &=< 0.108, 0.016 > + < 0.012, 0.064 > + < 0.072, 0.144 > + < 0.008, 0.576 > \\ &=< 0.2, 0.8 > \end{aligned}$$

Conditioning

		toothache	¬toothache	
		catch	¬catch	catch
cavity	catch	0.108	0.012	0.072
	¬catch	0.016	0.064	0.144

Using the product rule in $Pr(Y) = \sum_z Pr(Y, Z = z)$ we obtain the **conditioning** rule:
 $Pr(Y) = \sum_z Pr(Y|z)Pr(z)$

To get the conditional probabilities, we use the definition:

$$\begin{aligned} Pr(cavity|toothache) &= \frac{Pr(cavity \wedge toothache)}{Pr(toothache)} \\ &= \frac{0.108 + 0.012}{0.108 + 0.012 + 0.016 + 0.064} = 0.6 \end{aligned}$$

$$\begin{aligned} Pr(\neg cavity|toothache) &= \frac{Pr(\neg cavity \wedge toothache)}{Pr(toothache)} \\ &= \frac{0.016 + 0.064}{0.108 + 0.012 + 0.016 + 0.064} = 0.4 \end{aligned}$$

Normalization

Notice that $Pr(\text{toothache})$ appears in denominator in both equations in the preceding conditional probability calculations. We can view it as a **normalization constant**, denoted α , that ensures that the conditional probability distribution $Pr(\text{Cavity}|\text{toothache})$ sums to 1.

Then we can write:

$$\begin{aligned} Pr(\text{Cavity}|\text{toothache}) &= \alpha Pr(\text{Cavity}, \text{toothache}) \\ &= \alpha [Pr(\text{Cavity}, \text{toothache}, \text{catch}) + Pr(\text{Cavity}, \text{toothache}, \neg\text{catch})] \\ &= \alpha [< 0.108, 0.16 > + < 0.012, 0.064 >] \\ &= \alpha < 0.12, 0.08 > = < 0.6, 0.4 > \end{aligned}$$

All of the preceding can be summarized in a general inference procedure using full joint probability distributions. Let E be a list of evidence variables, e be a list of observed values for the evidence variables, and Y be the remaining unobserved variables. Then:

$$Pr(X|e) = \alpha Pr(X, e) = \alpha \sum_y Pr(X, e, y)$$

Great, so we're done! All we need is a full joint probability distribution and we can answer any query. Unfortunately, not practical.

- ▶ For a domain of n boolean variables, we have a table of size $O(2^n)$

So we need different approaches, which we cover next.

Independence

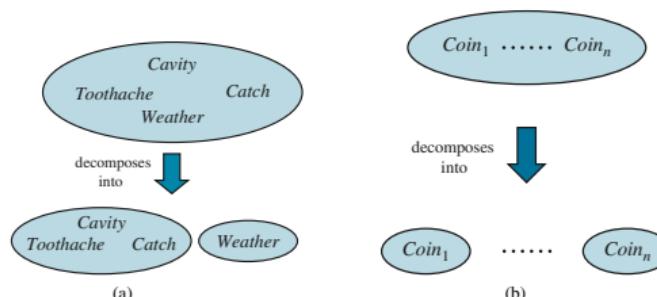
Weather is not affected by teeth. So we can assert:

$$Pr(\text{cloud}|\text{toothache}, \text{catch}, \text{cavity}) = Pr(\text{cloud})$$

In general, if a and b are independent:

$$P(a|b) = P(a) \text{ or } P(b|a) = P(b) \text{ or } P(a \wedge b) = P(a)P(b)$$

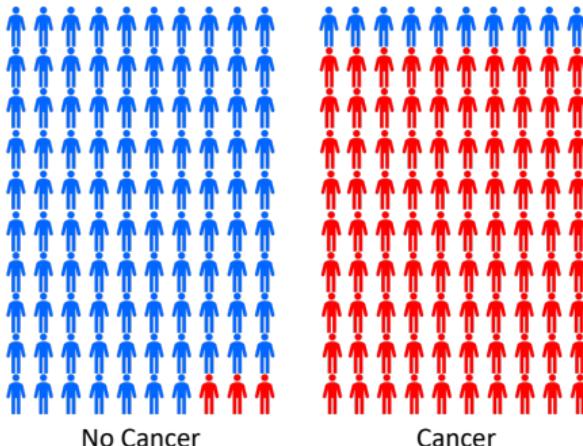
Can be a big help. For example, if you have n independent coin flips, then instead of 2^n full joint table, you have a product of n distributions $Pr(C_i)$.



Unfortunately, independence rarely holds in the real world.

A Medical Screening Example

A cancer with occurrence rate of 1% (.01) has a “90% accurate” test, and:



False positive rate: .03, False negative rate: 0.10

Questions:

- ▶ If we screen someone, what is the probability that they test positive?
- ▶ If someone tests positive, what is the probability that they have cancer?

Bayes' Theorem

Using the symmetry $p(X, Y) = p(Y, X)$ and the product rule:

$$\begin{aligned} p(X, Y) &= p(Y, X) \\ p(Y|X)p(X) &= p(X|Y)p(Y) \\ p(Y|X) &= \frac{p(X|Y)p(Y)}{p(X)} \end{aligned}$$

where the denominator $p(X)$ is a normalizing constant:

$$p(X) = \sum p(X|Y)p(Y)$$

This is called *Bayes' Theorem* or *Bayes' Rule*.

We use Bayes' Theorem to update our beliefs after observing evidence. For example:

- ▶ Before we run the test, the *prior probability* that someone has cancer is $p(C)$
- ▶ After we run the test, we use Bayes' Theorem to calculate the *posterior probability* $p(C|T)$

The *posterior probability* is our new belief after a Bayesian update.

Analysis of Medical Screening Example

With our probabilistic machinery we can now analyze our cancer screening example. First, we model the problem in the language of Bayesian probability theory:

$$p(C = 1) = \frac{1}{100} \quad (\text{Prior probability that someone has cancer})$$

$$p(C = 0) = \frac{99}{100} \quad (\text{Prior probability that someone has no cancer})$$

$$p(T = 1|C = 1) = \frac{90}{100} \quad (\text{Conditional probability of positive test given cancer})$$

$$p(T = 0|C = 1) = \frac{10}{100} \quad (\text{Conditional probability of negative test given cancer})$$

$$p(T = 1|C = 0) = \frac{3}{100} \quad (\text{Conditional probability of positive test given no cancer})$$

$$p(T = 0|C = 0) = \frac{97}{100} \quad (\text{Conditional probability of negative test given no cancer})$$

Now we can answer the two questions we posed at the outset:

- ▶ If we screen someone, what is the probability that they test positive?
- ▶ If someone tests positive, what is the probability that they have cancer?

Analysis of Medical Screening Example

If we screen someone, probability that they test positive:

$$p(C = 1) = \frac{1}{100}$$

$$p(C = 0) = \frac{99}{100}$$

$$p(T = 1|C = 1) = \frac{90}{100}$$

$$p(T = 0|C = 1) = \frac{10}{100}$$

$$p(T = 1|C = 0) = \frac{3}{100}$$

$$p(T = 0|C = 0) = \frac{97}{100}$$

$$\begin{aligned} p(T = 1) &= p(T = 1|C = 0)p(C = 0) + p(T = 1|C = 1)p(C = 1) \\ &= \frac{3}{100} \times \frac{99}{100} + \frac{90}{100} \times \frac{1}{100} \\ &= \frac{387}{10,000} \\ &= .0387 \end{aligned}$$

If someone tests positive, probability they have cancer:

$$\begin{aligned} p(C = 1|T = 1) &= \frac{p(T = 1|C = 1)p(C = 1)}{p(T = 1)} \\ &= \frac{90}{100} \times \frac{1}{100} \times \frac{10,000}{387} \\ &= \frac{90}{387} \\ &\approx 0.23 \end{aligned}$$