

# Relational Algebra

# Relational Algebra

Relational model specifies structures and constraints, relational algebra provides retrieval operations

- ▶ Formal foundation for relational model operations
- ▶ Basis for internal query optimization in RDBMS
- ▶ Parts of relational algebra found in SQL

## Basic Rules

- ▶ Relational algebra expressions operate on relations and produce relations as results
- ▶ Relational algebra expressions can be chained

## SELECT

$$\sigma_{<\text{condition}>} (R)$$

- ▶  $R$  is the name of a relation
- ▶  $<\text{condition}>$  is a boolean condition on the values of attributes in the tuples of  $R$

The select operation returns all the tuples from  $R$  for which  $<\text{condition}>$  is true.

## SELECT Example

Given the following data for pet:

shelter_id	id	name	breed
1	1	Chloe	Mix
1	2	Dante	GSD
1	3	Heidi	Dachshund
2	1	Bailey	Mix
2	2	Sophie	Lab
2	3	Heidi	Dachshund

$\sigma_{breed='mix'}(pet)$  returns:

shelter_id	id	name	breed
1	1	Chloe	Mix
2	1	Bailey	Mix

# Properties of SELECT

- ▶ Result of  $\sigma_{<condition>}(R)$  has same schema as  $R$ , i.e., same attributes
- ▶ SELECT is commutative, e.g.,  
$$\sigma_{<c1>}(\sigma_{<c2>}(R)) = \sigma_{<c2>}(\sigma_{<c1>}(R))$$
- ▶ Cascaded SELECTs can be replaced by single SELECT with conjunction of conditions, e.g.  
$$\sigma_{<c1>}(\sigma_{<c2>}(R)) = \sigma_{<c1>AND<c2>}(R)$$
- ▶ Result of  $\sigma_{<condition>}(R)$  has equal or fewer tuples than  $R$

# PROJECT

$$\pi_{<\text{attributelist}>} (R)$$

- ▶  $R$  is the name of a relation
- ▶  $<\text{attributelist}>$  is a subset of the attributes of relation  $R$

The project operation returns all the tuples in  $R$  but with only the attributes in  $<\text{attribute-list}>$

## PROJECT Example

shelter_id	id	name	breed
1	1	Chloe	Mix
1	2	Dante	GSD
1	3	Heidi	Dachshund
2	1	Bailey	Mix
2	2	Sophie	Lab
2	3	Heidi	Dachshund

$$\pi_{name, breed}(pet) =$$

name	breed
Chloe	Mix
Dante	GSD
Heidi	Dachshund
Bailey	Mix
Sophie	Lab

Notice that the duplicate tuple <Heidi, Dachshund> was removed. **Georgia Tech**  
Results of relational algebra operations are sets.

## Properties of PROJECT

- ▶ Number of tuples returned by PROJECT is less than or equal to the number of tuples in the input relation because result is a set, i.e.,  
 $|\pi_{\langle attrs \rangle}(R)| \leq |R|$ 
  - ▶ What if  $\langle attrs \rangle$  includes a key of  $R$ ?
- ▶ PROJECT is **not** commutative. In fact  $\pi_{\langle attrs1 \rangle}(\pi_{\langle attrs2 \rangle}(R))$  is only a correct expression if  $\langle attrs2 \rangle$  contains the attributes in  $\langle attrs1 \rangle$ . In this case the result is simply  $\pi_{\langle attrs1 \rangle}(R)$ .

## Combining PROJECT and SELECT

shelter_id	id	name	breed
1	1	Chloe	Mix
1	2	Dante	GSD
1	3	Heidi	Dachshund
2	1	Bailey	Mix
2	2	Sophie	Lab
2	3	Heidi	Dachshund

$\pi_{name}(\sigma_{breed='Mix'}(pet))$  produces:

name
Chloe
Bailey

## Intermediate Results

Previous **in-line expression** could be split up into multiple steps with named intermediate results.

$$\pi_{name}(\sigma_{breed='Mix'}(pet))$$

becomes:

$$\begin{aligned}MIXES &\leftarrow \sigma_{breed='Mix'}(pet) \\RESULT &\leftarrow \pi_{name}(MIXES)\end{aligned}$$

# RENAME

- ▶ Rename relation  $R$  to  $S$ :

$$\rho_S(R)$$

- ▶ Rename attributes of  $R$  to  $B_1, \dots, B_n$ :

$$\rho_{(B_1, \dots, B_n)}(R)$$

- ▶ Rename  $R$  to  $S$  and attributes to  $B_1, \dots, B_n$ :

$$\rho_{S(B_1, \dots, B_n)}(R)$$

# Binary Operators

- ▶ UNION,  $R \cup S$ , is set of all tuples in either  $R$  or  $S$
- ▶ INTERSECTION,  $R \cap S$ , is set of all tuples in both  $R$  and  $S$
- ▶ SET DIFFERENCE,  $R - S$ , is set of all tuples in  $R$  but not in  $S$

Operands must be **union compatible**, or **type compatible**. For  $R$  and  $S$  to be union compatible:

- ▶ Degree of  $R$  must be same as degree of  $S$
- ▶ For each attribute  $A_i$  in  $R$  and  $B_i$  in  $S$ ,  $\text{dom}(A_i) = \text{dom}(B_i)$

## Cartesian Product

$R \times S$  Creates "super-tuples" by concatenating every tuple in  $R$  with every tuple in  $S$ .

$$R(A_1, \dots, A_n) \times S(B_1, \dots, B_m) = Q(A_1, \dots, A_n, B_1, \dots, B_m)$$

Notice that

- ▶  $Q$  has degree  $n + m$
- ▶  $|q(Q)| = |r(R)| \times |s(S)|$

Note that the book abuses notation a bit and writes that last bullet as  
 $|Q| = |R| \times |S|$

# Cartesian Product Example

shelter

id	name
1	Howell
2	Mansell

pet

shelter_id	id	name	breed
1	1	Chloe	Mix
1	2	Dante	GSD
1	3	Heidi	Dachshund
2	1	Bailey	Mix
2	2	Sophie	Lab
2	3	Heidi	Dachshund

## Cross Product Example

sid	sname	shelter_id	pid	pname	breed
1	Howell	1	1	Chloe	Mix
2	Mansell	1	1	Chloe	Mix
1	Howell	1	2	Dante	GSD
2	Mansell	1	2	Dante	GSD
1	Howell	1	3	Heidi	Dachshund
2	Mansell	1	3	Heidi	Dachshund
1	Howell	2	1	Bailey	Mix
2	Mansell	2	1	Bailey	Mix
1	Howell	2	2	Sophie	Lab
2	Mansell	2	2	Sophie	Lab
1	Howell	2	3	Heidi	Dachshund
2	Mansell	2	3	Heidi	Dachshund

Note that we've also done a RENAME to disambiguate name and id:

$\rho(sid, sname, shelter\_id, pid, pname, breed)(shelter \times pet)$

# Cross Product and Select

Cross product meaningful when combined with SELECT.

$$\sigma_{sid=shelter\_id}(\rho_{(sid,sname,shelter\_id,pid,pname,breed)}(shelter \times pet))$$

sid	sname	shelter_id	pid	pname	breed
1	Howell	1	1	Chloe	Mix
1	Howell	1	2	Dante	GSD
1	Howell	1	3	Heidi	Dachshund
2	Mansell	2	1	Bailey	Mix
2	Mansell	2	2	Sophie	Lab
2	Mansell	2	3	Heidi	Dachshund

*CROSSED*  $\leftarrow$  shelter  $\times$  pet

*RENAMED*  $\leftarrow$   $\rho_{(sid,sname,shelter\_id,pid,pname,breed)}(CROSSED)$

*RESULT*  $\leftarrow$   $\sigma_{sid=shelter\_id}(RENAMED)$

# Join

JOIN is a CARTESIAN PRODUCT followed by SELECT

$$R \bowtie_{\langle \text{joincondition} \rangle} S$$

Where

- ▶  $R$  and  $S$  are relations
- ▶  $\langle \text{joincondition} \rangle$  is a boolean condition on values of tuples from  $R$  and  $S$

$R \bowtie_{\langle \text{joincondition} \rangle} S$  returns the tuples in  $R \times S$  that satisfy the  $\langle \text{joincondition} \rangle$

## Join Conditions

$\langle joincondition \rangle$  is of the form  $A_i\theta B_j$

- ▶  $A_i$  is an attribute of  $R$ ,  $B_j$  is an attribute of  $S$
- ▶  $dom(A_i) = dom(B_j)$
- ▶  $\theta$  is one of  $\{=, \neq, <>, <, \leq, >, \geq\}$

A  $\langle joincondition \rangle$  can be a conjunction of simple conditions, e.g.:

$\langle c_1 \rangle AND \langle c_2 \rangle ... AND \langle c_n \rangle$

# Join Example

worker

id	name	supervisor_id	shelter_id
1	Tom	NULL	1
2	Jie	1	1
3	Ravi	2	1
4	Alice	2	1
5	Aparna	NULL	2
6	Bob	5	2
7	Xaoxi	6	2
8	Rohan	6	2

shelter

id	name
1	Howell
2	Mansell

## Join Example

$\text{worker} \bowtie_{\text{shelter\_id}=\text{sid}} \rho_{(\text{sid}, \text{pname})}(\text{shelter})$

id	name	supervisor_id	shelter_id	sid	sname
1	Tom	NULL	1	1	Howell
2	Jie	1	1	1	Howell
3	Ravi	2	1	1	Howell
4	Alice	2	1	1	Howell
5	Aparna	NULL	2	2	Mansell
6	Bob	5	2	2	Mansell
7	Xaoxi	6	2	2	Mansell
8	Rohan	6	2	2	Mansell

Notice that we had to use renaming of attributes in *shelter*.

A join operation in which the comparison operator  $\theta$  is  $=$  is called an **equijoin**.

## Natural Join

Notice that the `shelter_id` attribute was repeated in the previous equijoin result. A NATURAL JOIN is an equijoin in which the redundant attribute has been removed.

$$R * S$$

Where

- ▶  $R$  and  $S$  have an attribute with the same name and same domain which is automatically chosen as the equijoin attribute

## Natural Join Example

Recall the first join example. If we rename the `id` attribute to `shelter_id` we can use a natural join:

$$\rho_{(shelter\_id, sname)}(shelter) * worker$$

shelter_id	sname	id	name	supervisor_id
1	Howell	1	Tom	NULL
1	Howell	2	Jie	1
1	Howell	3	Ravi	2
1	Howell	4	Alice	2
2	Mansell	5	Aparna	NULL
2	Mansell	6	Bob	5
2	Mansell	7	Xaoxi	6
2	Mansell	8	Rohan	6

## Outer Joins

The joins we've discussed so far have been inner joins. Result relations of inner joins include only tuples from the joined tables that match the join condition.

Outer join results include tuples that matched, and tuples that didn't match the join condition.

## Left Outer Join

$$R \bowtie_{\langle \text{joincondition} \rangle} S$$

Where

- ▶  $R$  and  $S$  are relations
- ▶  $\langle \text{joincondition} \rangle$  is a boolean condition on values of tuples from  $R$  and  $S$

$R \bowtie_{\langle \text{joincondition} \rangle} S$  returns the tuples in  $R \times S$  that satisfy the  $\langle \text{joincondition} \rangle$  as well as the tuples from  $R$  that don't match the join condition. In the result relation the unmatched tuples from  $R$  are null-padded to give them the correct degree in the result.

# Left Outer Join Example

author

author_id	first_name	last_name
1	John	McCarthy
2	Dennis	Ritchie
3	Ken	Thompson
4	Claude	Shannon
5	Alan	Turing
6	Alonzo	Church
7	Perry	White
8	Moshe	Vardi
9	Roy	Batty

book

book_id	book_title	month	year	editor
1	CACM	April	1960	8
2	CACM	July	1974	8
3	BST	July	1948	2
4	LMS	November	1936	7
5	Mind	October	1950	NULL
6	AMS	Month	1941	NULL
7	AAAI	July	2012	9

# Authors and Edited Books

Show all the authors. For all the authors who are editors, show their books.

$R \bowtie_{author\_id=editor} S$

author_id	first_name	last_name	book_id	book_title	month	year	editor
8	Moshe	Vardi	1	CACM	April	1960	8
8	Moshe	Vardi	2	CACM	July	1974	8
2	Dennis	Ritchie	3	BST	July	1948	2
7	Perry	White	4	LMS	November	1936	7
9	Roy	Batty	7	AAAI	July	2012	9
9	Roy	Batty	8	NIPS	July	2012	9
1	John	McCarthy	NULL	NULL	NULL	NULL	NULL
3	Ken	Thompson	NULL	NULL	NULL	NULL	NULL
4	Claude	Shannon	NULL	NULL	NULL	NULL	NULL
5	Alan	Turing	NULL	NULL	NULL	NULL	NULL
6	Alonzo	Church	NULL	NULL	NULL	NULL	NULL

Notice how attribute values are padded to the right in a left outer join.

# Right Outer Join

$$R \bowtie_{\langle \text{joincondition} \rangle} S$$

Where

- ▶  $R$  and  $S$  are relations
- ▶  $\langle \text{joincondition} \rangle$  is a boolean condition on values of tuples from  $R$  and  $S$

$R \bowtie_{\langle \text{joincondition} \rangle} S$  returns the tuples in  $R \times S$  that satisfy the  $\langle \text{joincondition} \rangle$  as well as the tuples from  $S$  that don't match the join condition. In the result relation the unmatched tuples from  $S$  are null-padded to give them the correct degree in the result.

# Right Outer Join Example

Show all the books. For books with editors, show their editors.

$$R \bowtie_{\text{author\_id}=\text{editor}} S$$

author_id	first_name	last_name	book_id	book_title	month	year	editor
8	Moshe	Vardi	1	CACM	April	1960	8
8	Moshe	Vardi	2	CACM	July	1974	8
2	Dennis	Ritchie	3	BST	July	1948	2
7	Perry	White	4	LMS	November	1936	7
NULL	NULL	NULL	5	Mind	October	1950	NULL
NULL	NULL	NULL	6	AMS	Month	1941	NULL
9	Roy	Batty	7	AAAI	July	2012	9
9	Roy	Batty	8	NIPS	July	2012	9

Notice how attribute values are padded to the left in a right outer join.

## Full Outer Join

$$R \bowtie_{\langle \text{joincondition} \rangle} S$$

Where

- ▶  $R$  and  $S$  are relations
- ▶  $\langle \text{joincondition} \rangle$  is a boolean condition on values of tuples from  $R$  and  $S$

$R \bowtie_{\langle \text{joincondition} \rangle} S$  returns the tuples in  $R \times S$  that satisfy the  $\langle \text{joincondition} \rangle$  as well as the tuples from both  $R$  and  $S$  that don't match the join condition. In the result relation the unmatched tuples are null-padded to give them the correct degree in the result.

# Full Outer Join Example

Show all authors and books, matching editors with their books.

$R \bowtie_{author\_id=editor} S$

author_id	first_name	last_name	book_id	book_title	month	year	editor
8	Moshe	Vardi	1	CACM	April	1960	8
8	Moshe	Vardi	2	CACM	July	1974	8
2	Dennis	Ritchie	3	BST	July	1948	2
7	Perry	White	4	LMS	November	1936	7
9	Roy	Batty	7	AAAI	July	2012	9
9	Roy	Batty	8	NIPS	July	2012	9
1	John	McCarthy	NULL	NULL	NULL	NULL	NULL
3	Ken	Thompson	NULL	NULL	NULL	NULL	NULL
4	Claude	Shannon	NULL	NULL	NULL	NULL	NULL
5	Alan	Turing	NULL	NULL	NULL	NULL	NULL
6	Alonzo	Church	NULL	NULL	NULL	NULL	NULL
NULL	NULL	NULL	5	Mind	October	1950	NULL
NULL	NULL	NULL	6	AMS	Month	1941	NULL

## Review Question 1

Given the  $r(book)$ :

book_id	book_title	month	year	editor
1	CACM	April	1960	8
2	CACM	July	1974	8
3	BST	July	1948	2
4	LMS	November	1936	7
5	Mind	October	1950	7
6	AMS	Month	1941	7
7	AAAI	July	2012	9
8	NIPS	July	2012	9

How many tuples are in  $\pi_{book\_title}(book)$ ?

## Review Question 1 Answer

7:

book_title
CACM
BST
LMS
Mind
AMS
AAAI
NIPS

The *book\_title* appears twice in *book* and the result of a relational algebra expression is a set.

## Review Question 2

Given the relation  $r(book)$ :

book_id	book_title	month	year	editor
1	CACM	April	1960	8
2	CACM	July	1974	8
3	BST	July	1948	2
4	LMS	November	1936	7
5	Mind	October	1950	7
6	AMS	Month	1941	7
7	AAAI	July	2012	9
8	NIPS	July	2012	9

Which books were published before 1960 or after 2000?

## Review Question 2

Which books were published before 1960 or after 2000?

$$\sigma_{\text{year} < 1960}(\text{book}) \cup \sigma_{\text{year} > 2000}(\text{book})$$

book_id	book_title	month	year	editor
3	BST	July	1948	2
4	LMS	November	1936	7
5	Mind	October	1950	7
6	AMS	Month	1941	7
7	AAAI	July	2012	9
8	NIPS	July	2012	9

## Review Question 3

Given:

worker

id	name	supervisor_id	shelter_id
1	Tom	NULL	1
2	Jie	1	1
3	Ravi	2	1
4	Alice	2	1
5	Aparna	NULL	2
6	Bob	5	2
7	Xaoxi	6	2
8	Rohan	6	2

shelter

id	name
1	Howell
2	Mansell

How would we find the names of all the workers who work at Mansell?

## Review Question 3 Answer

How would we find all the workers who work at Mansell?

$SHELTERS \leftarrow \rho_{sid, sname}(shelter)$

$MANSELLERS \leftarrow$

$worker \bowtie_{shelter\_id = sid \wedge sname = 'Mansell'} (SHELTERS)$

Gives:

id	name	supervisor_id	shelter_id	sid	sname
5	Aparna	NULL	2	2	Mansell
6	Bob	5	2	2	Mansell
7	Xaoxi	6	2	2	Mansell
8	Rohan	6	2	2	Mansell

then ...

## Review Question 3 Answer

$\pi_{name}(MANSELLERS)$   
gives:

name
Aparna
Bob
Xaoxi
Rohan

Full inline expression:

$$\pi_{name}(\text{worker} \bowtie_{shelter\_id=sid \wedge sname='Mansell'} \rho_{(sid,sname)}(shelter))$$