

Loss Functions

CS 4277 Deep Learning

Kennesaw State University

Loss Functions

The goal of a (eager) machine learning algorithm is to find the parameters of a model that produces the best possible mapping from inputs to outputs.

- ▶ We do this using a training data set $\{x_i, y_i\}_{i=1}^N$
- ▶ Training uses feedback from the mismatch between the model's predicted \hat{y} s and the ground truth y s.
- ▶ A *loss function* returns a single number that represents this mismatch.
- ▶ So finding the best possible mapping from inputs to outputs reduces to minimizing the loss function.

Density Estimation

Say we have N observations of a scalar x which we denote with $\mathbf{x} = (x_1, \dots, x_N)$.

- ▶ Estimating the distribution given the data is known as *density estimation*.
- ▶ We must assume a distribution, so we're estimating the parameters of the distribution.
- ▶ We assume data points are drawn independently and are identically-distributed.
 - ▶ This is known as the i.i.d. assumption

Example: Likelihood of the Gaussian

Since our data set is i.i.d., the probability of the data set is

$$p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \mathcal{N}(x_n|\mu, \sigma^2)$$

This is known as the *likelihood function* for the Gaussian.

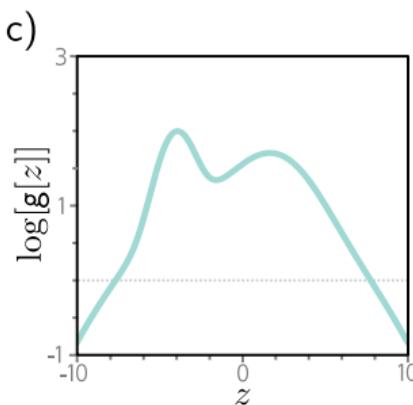
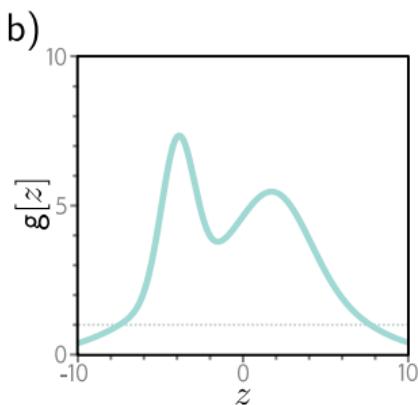
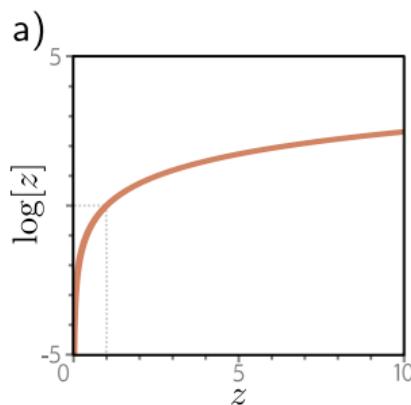
Maximum Likelihood

Finding the parameters of a distribution that maximize the probability of the observed data is known as *maximum likelihood estimation* (MLE).

In practice, we transform likelihood functions into log likelihood functions.

Why log:

- ▶ Log of a function monotonically increasing and concave – $\operatorname{argmax} \ln(f) = \operatorname{argmax} f$
- ▶ Log easy to work with: $\ln(ab) = \ln(a) + \ln(b)$, $\ln\left(\frac{a}{b}\right) = \ln(a) - \ln(b)$, $\ln e^x = x$
- ▶ Multiplying probabilities can underflow – summing logs avoids this problem



Log Likelihood of Gaussian

For the Gaussian likelihood function we saw earlier, the log likelihood

$$p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \mathcal{N}(x_n|\mu, \sigma^2)$$

$$p(\mathbf{x}|\mu, \sigma^2) = \prod_{n=1}^N \frac{1}{(2\pi\sigma^2)^{\frac{1}{2}}} \exp\left(-\frac{1}{2\sigma^2}(x_n - \mu)^2\right)$$

$$\ln p(\mathbf{x}|\mu, \sigma^2) = \sum_{n=1}^N \ln\left(\frac{1}{(2\pi\sigma^2)^{\frac{1}{2}}} \exp\left(-\frac{1}{2\sigma^2}(x_n - \mu)^2\right)\right)$$

$$\ln p(\mathbf{x}|\mu, \sigma^2) = \sum_{n=1}^N \ln(1) - \ln(\sqrt{2\pi}) - \ln(\sigma) - \frac{(x_i - \mu)^2}{2\sigma^2}$$

Maximum Likelihood of Gaussian

If we take the partial derivative of the Gaussian log likelihood function with respect to μ , set it to zero, and solve for μ we get:

$$\mu_{ML} = \frac{1}{N} \sum_{n=1}^N x_n$$

If we take the partial derivative with respect to σ^2 we get:

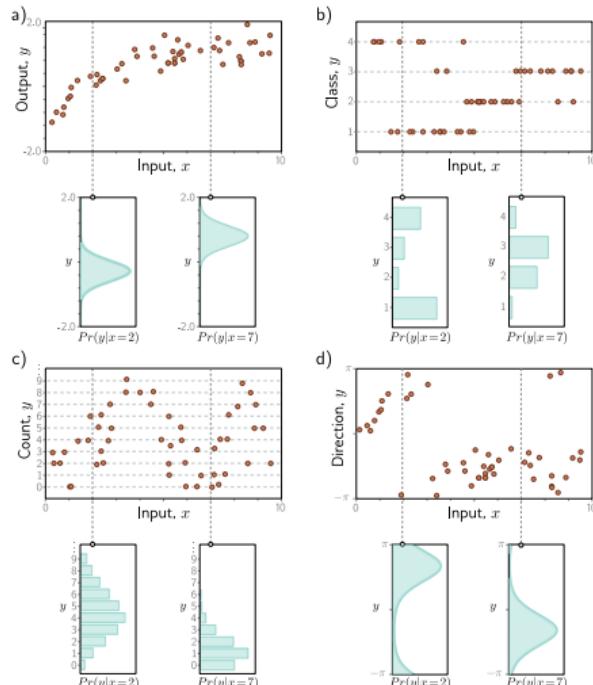
$$\sigma_{ML}^2 = \frac{1}{N} \sum_{n=1}^N (x_n - \mu_{ML})^2$$

These should look familiar. They are the sample mean and sample variance of the Gaussian.

Loss Functions and Machine Learning Models

We seek a model $f(\mathbf{x}, (\phi))$ that computes a \hat{y} given an \mathbf{x} .

- We can recast this problem as the computation of a conditional probability $p(y_i | \mathbf{x}_i)$.
- Minimizing the loss corresponds to maximizing this conditional probability.



General Maximum Likelihood Criterion

We choose a parametric distribution defined over the output domain \mathbf{y} then train our model to compute the parameters, θ of this distribution.

- If we choose a Gaussian distribution, then $\theta = \{\mu, \sigma^2\}$.

We want to find the parameters of the model $\hat{\phi}$ that maximize the conditional probability distribution $P(\mathbf{y}_i|\theta_i)$ for all \mathbf{y}_i s.

$$\hat{\phi} = \operatorname{argmax}_{\phi} \left(\prod_{i=1}^I p(\mathbf{y}_i | \mathbf{x}_i) \right)$$

$$\hat{\phi} = \operatorname{argmax}_{\phi} \left(\prod_{i=1}^I p(\mathbf{y}_i | \theta_i) \right)$$

$$\hat{\phi} = \operatorname{argmax}_{\phi} \left(\prod_{i=1}^I p(\mathbf{y}_i | f(\mathbf{x}, \phi)) \right)$$

Maximizing Log-Likelihood

Recalling that the total likelihood of the training data is:

$$P(\mathbf{y}_1, \dots, \mathbf{y}_I | \mathbf{x}_1, \dots, \mathbf{x}_I) = \prod_{i=1}^I p(\mathbf{y}_i | \mathbf{x}_i)$$

which is impractical due to underflow, so we prefer to maximize the log-likelihood:

$$\hat{\phi} = \underset{\phi}{\operatorname{argmax}} \left(\prod_{i=1}^I p(\mathbf{y}_i | \mathbf{f}(\mathbf{x}, \phi)) \right)$$

$$\hat{\phi} = \underset{\phi}{\operatorname{argmax}} \left(\log \left(\prod_{i=1}^I p(\mathbf{y}_i | \mathbf{f}(\mathbf{x}, \phi)) \right) \right)$$

$$\hat{\phi} = \underset{\phi}{\operatorname{argmax}} \left(\sum_{i=1}^I \log(p(\mathbf{y}_i | \mathbf{f}(\mathbf{x}, \phi))) \right)$$

Minimizing Negative Log-Likelihood

By convention we minimize the loss function. We can turn a maximization problem into a minimization problem by multiplying by -1.

$$\hat{\phi} = \operatorname{argmin}_{\phi} \left(- \sum_{i=1}^I \log(p(\mathbf{y}_i | \mathbf{f}(\mathbf{x}, \phi)) \right)$$
$$\hat{\phi} = \operatorname{argmin}_{\phi} (L(\phi))$$

Which is the final form of our loss function $L(\phi)$

Inference

Our network now computes a probability distribution over \mathbf{y} instead of predicting \hat{y} . To get a prediction, we return the maximum of the distribution:

$$\hat{\mathbf{y}} = \operatorname{argmax}_{\mathbf{y}}(p(\mathbf{y}|\mathbf{f}(\mathbf{x}, \phi)))$$

This is often computed in terms of the parameters θ predicted by the model. E.g., for Gaussian the maximum is at μ

Recipe for Constructing and Using Loss Functions

Now that we understand MLE for loss functions, we can create a recipe for constructing loss functions for training data $\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^I$ using the maximum likelihood approach:

1. Choose a suitable probability distribution $P(\mathbf{y}|\boldsymbol{\theta})$ defined over the predictions (output domain) \mathbf{y} with distribution parameters $\boldsymbol{\theta}$.
2. Set the machine learning model $\mathbf{f}(\mathbf{x}, \boldsymbol{\phi})$ to predict one or more of these parameters, so $\boldsymbol{\theta} = \mathbf{f}(\mathbf{x}, \boldsymbol{\phi})$ and $P(\mathbf{y}|\boldsymbol{\theta}) = P(\mathbf{y}|\mathbf{f}(\mathbf{x}, \boldsymbol{\phi}))$.
3. To train the model, find the network parameters $\hat{\boldsymbol{\phi}}$ that minimize the negative log-likelihood loss function over the training data pairs $\{\mathbf{x}_i, \mathbf{y}_i\}$:

$$\hat{\boldsymbol{\phi}} = \underset{\boldsymbol{\phi}}{\operatorname{argmin}}(L(\boldsymbol{\phi})) = \underset{\boldsymbol{\phi}}{\operatorname{argmin}}\left(-\sum_{i=1}^I \log(p(\mathbf{y}_i|\mathbf{f}(\mathbf{x}, \boldsymbol{\phi}))\right)$$

4. To perform inference for a new test example \mathbf{x} , return either the full distribution $P(\mathbf{y}|\mathbf{f}(\mathbf{x}, \boldsymbol{\phi}))$ or the value where the distribution is maximized.

Now we see this recipe applied to three common problems: univariate regression, binary classification, and multiclass classification.

Example 1: Univariate Regression

Goal: predict a single output $y \in \mathbb{R}$ from input \mathbf{x} using model $f(\mathbf{x}, \phi)$.

1. Choose a distribution over output y .

$$p(y|\mu, \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - \mu)^2}{2\sigma^2}\right)$$

2. Set the machine learning model $f(\mathbf{x}, \phi)$ to compute one or more parameters of the distribution. Here, we choose only the mean $\mu = f(\mathbf{x}, \phi)$:

$$p(y|f(\mathbf{x}, \phi), \sigma^2) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp\left(-\frac{(y - f(\mathbf{x}, \phi))^2}{2\sigma^2}\right)$$

3. Calculate the negative log-likelihood for our chosen parameter(s) to use as the loss function. We already did this for the Gaussian earlier; here we substitute $f(\mathbf{x}, \phi)$ for μ and y_i for x_i since we're calculating the distribution parameter over the output:

$$L(\phi) = \sum_{i=1}^I \ln(1) - \ln(\sqrt{2\pi}) - \ln(\sigma) - \frac{(y_i - f(\mathbf{x}, \phi))^2}{2\sigma^2}$$

3.1: Least-Squares Loss Function

Our negative log-likelihood function is unwieldy, so do some algebraic manipulations.

$$L(\phi) = \sum_{i=1}^I \ln(1) - \ln(\sqrt{2\pi}) - \ln(\sigma) - \frac{(y_i - f(\mathbf{x}, \phi))^2}{2\sigma^2}$$

First, remove terms that don't depend on ϕ :

$$L(\phi) = \sum_{i=1}^I \frac{(y_i - f(\mathbf{x}, \phi))^2}{2\sigma^2}$$

Next, remove denominator since it's a positive scaling factor that doesn't affect the position of the minimum:

$$L(\phi) = \sum_{i=1}^I (y_i - f(\mathbf{x}, \phi))^2$$

And this is the familiar least squares loss function you get when you assume that your outputs Y are i.i.d. and each $y_i \sim \mathcal{N}(y_i |, f(\mathbf{x}, \phi), \sigma^2)$

4: Inference

The model doesn't predict y directly, instead it predicts the mean of the Gaussian distribution assumed to generate y . We turn that into a predicted \hat{y} with:

$$\hat{y} = \underset{y}{\operatorname{argmax}}(p(y|f(\mathbf{x}, \hat{\phi}), \sigma^2))$$

Since for the univariate Gaussian the mean μ is the most likely value, inference is simply

$$\hat{y} = f(\mathbf{x}, \hat{\phi})$$

Heteroscedastic vs Homoscedatic Regression

- ▶ In *homoscedastic* regression, the variance is constant across all data points.
- ▶ In *heteroscedastic* regression, the variance varies as a function of the input data.

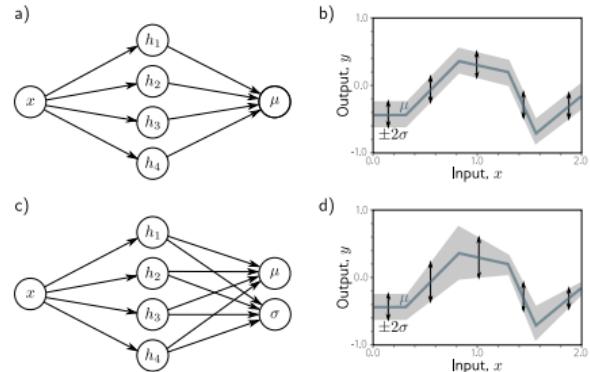
To handle heteroscedastic regression, you can train a network to compute both mean and variance.

$$\mu = f_1(\mathbf{x}, \phi)$$

$$\sigma^2 = f_2(\mathbf{x}, \phi)^2 \quad (\text{Square output to ensure positive})$$

The loss function is then:

$$\hat{\phi} = \underset{\phi}{\operatorname{argmin}} \left(- \sum_{i=1}^I \left(\log\left(\frac{1}{\sqrt{2\pi f_2(\mathbf{x}, \phi)^2}}\right) - \frac{(y_i - f_1(\mathbf{x}, \phi))^2}{2f_2(\mathbf{x}, \phi)^2} \right) \right)$$



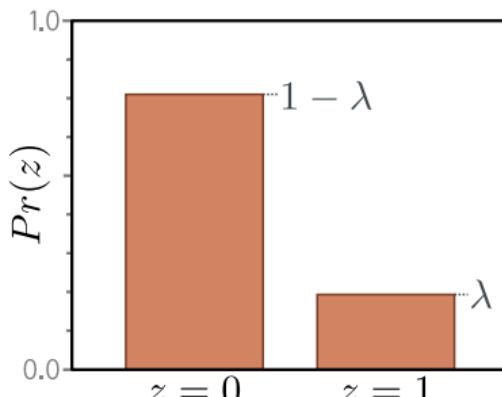
Example 2: Binary Classification

1. Choose Bernoulli distribution, which is defined over the output space $y \in \{0, 1\}$. The parameter λ represents the probability that $y = 1$.

$$p(y|\lambda) = \begin{cases} 1 - \lambda & y = 0, \\ \lambda & y = 1 \end{cases}$$

Or, equivalently:

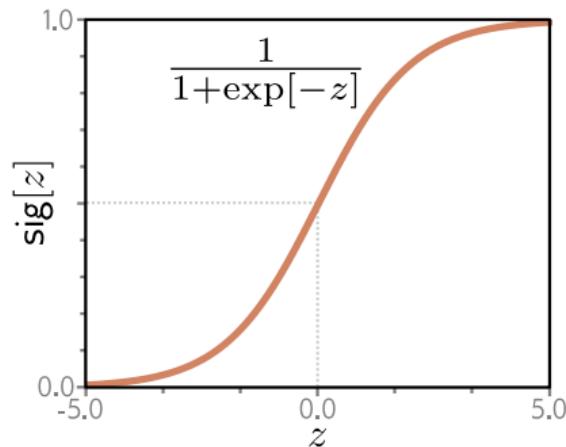
$$p(y|\lambda) = (1 - \lambda)^{1-y} \lambda^y$$



2: Predicting λ

- Set the network $f(\mathbf{x}, \phi)$ to predict the single parameter λ . Since λ is a probability, but the network can return a number outside $[0, 1]$, we squash the output to the required range using the *logistic sigmoid* function:

$$\text{sig}(z) = \frac{1}{1 + \exp(-z)}$$



Binary Loss Function and Inference

3. Calculate the negative log-likelihood for our parameter, λ , to use as the loss function.

The likelihood function is:

$$p(y|\mathbf{x}) = (1 - \text{sig}(f(\mathbf{x}, \phi)))^{1-y} \text{sig}(f(\mathbf{x}, \phi))^y$$

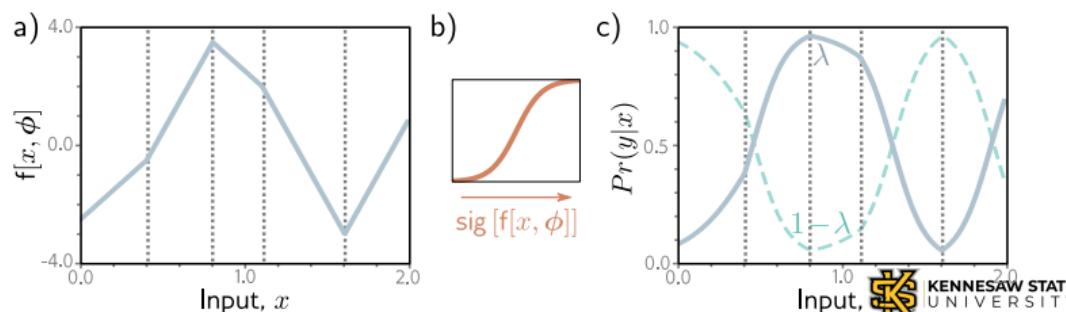
and the negative log-likelihood loss function is:

$$L(\phi) = \sum_{i=1}^I (1 - y_i) \log(1 - \text{sig}(f(\mathbf{x}_i, \phi))) - y_i \log(\text{sig}(f(\mathbf{x}_i, \phi)))$$

This is known as *binary cross-entropy loss*. (More on that later . . .)

4. Inference is simply:

$$\hat{y} = \begin{cases} 1 & \text{if } \lambda > 0.5, \\ 0 & \text{otherwise} \end{cases}$$



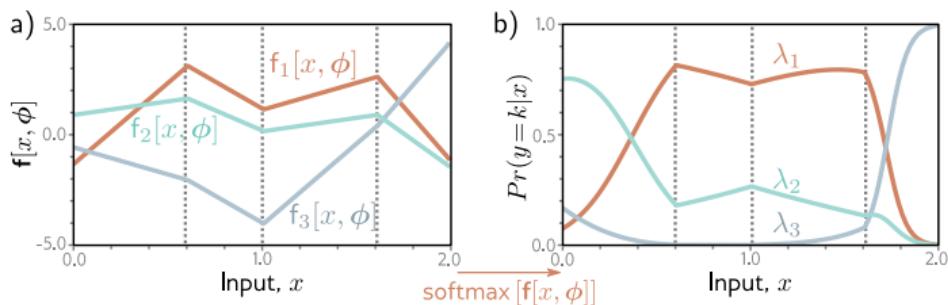
Example 3: Multiclass Classification

1. We choose a *categorical distribution* with K parameters $\lambda_1, \lambda_2, \dots, \lambda_K$ representing the probability of each of the corresponding K categories $y \in \{1, 2, \dots, K\}$.

$$p(y = k) = \lambda_k$$

2. We set the network $f(\mathbf{x}, \phi)$ to predict each of the K parameters for a given input \mathbf{x} . Since the parameters must sum to 1 to be a valid probability distribution and the network can produce arbitrary values, we pass each output through the *softmax* function:

$$\text{softmax}_k(z) = \frac{\exp(z_k)}{\sum_{k'=1}^K \exp(z_{k'})}$$



Multiclass Loss and Inference

3. Calculate the negative log-likelihood for our parameters, λ_k , to use as the loss function.

The likelihood function is:

$$p(y = k | \mathbf{x}) = \text{softmax}_k(f(\mathbf{x}, \phi))$$

The negative log-likelihood loss function is:

$$\begin{aligned} L(\phi) &= - \sum_{i=1}^I \log (\text{softmax}_{y_i}(f(\mathbf{x}_i, \phi))) \\ &= - \sum_{i=1}^I \left(f_{y_i}(\mathbf{x}_i, \phi) - \log \left(\sum_{k'=1}^K \exp(f_{k'}(\mathbf{x}_i, \phi)) \right) \right) \end{aligned}$$

where $f_{y_i}(\mathbf{x}_i, \phi)$ is the y_i th output and $f_{k'}(\mathbf{x}_i, \phi)$ is the k' th output of the network.

4. Inference is simply the most probable category:

$$\hat{y} = \operatorname{argmax}_k \left(p(y = k | f(\mathbf{x}, \hat{\phi})) \right)$$